

Deep Learning methods for pattern recognition

Ahmed Hammad

High Energy Accelerator Research Organization (KEK), Japan.

4th summer school at CTP, BUE

Fourth lecture

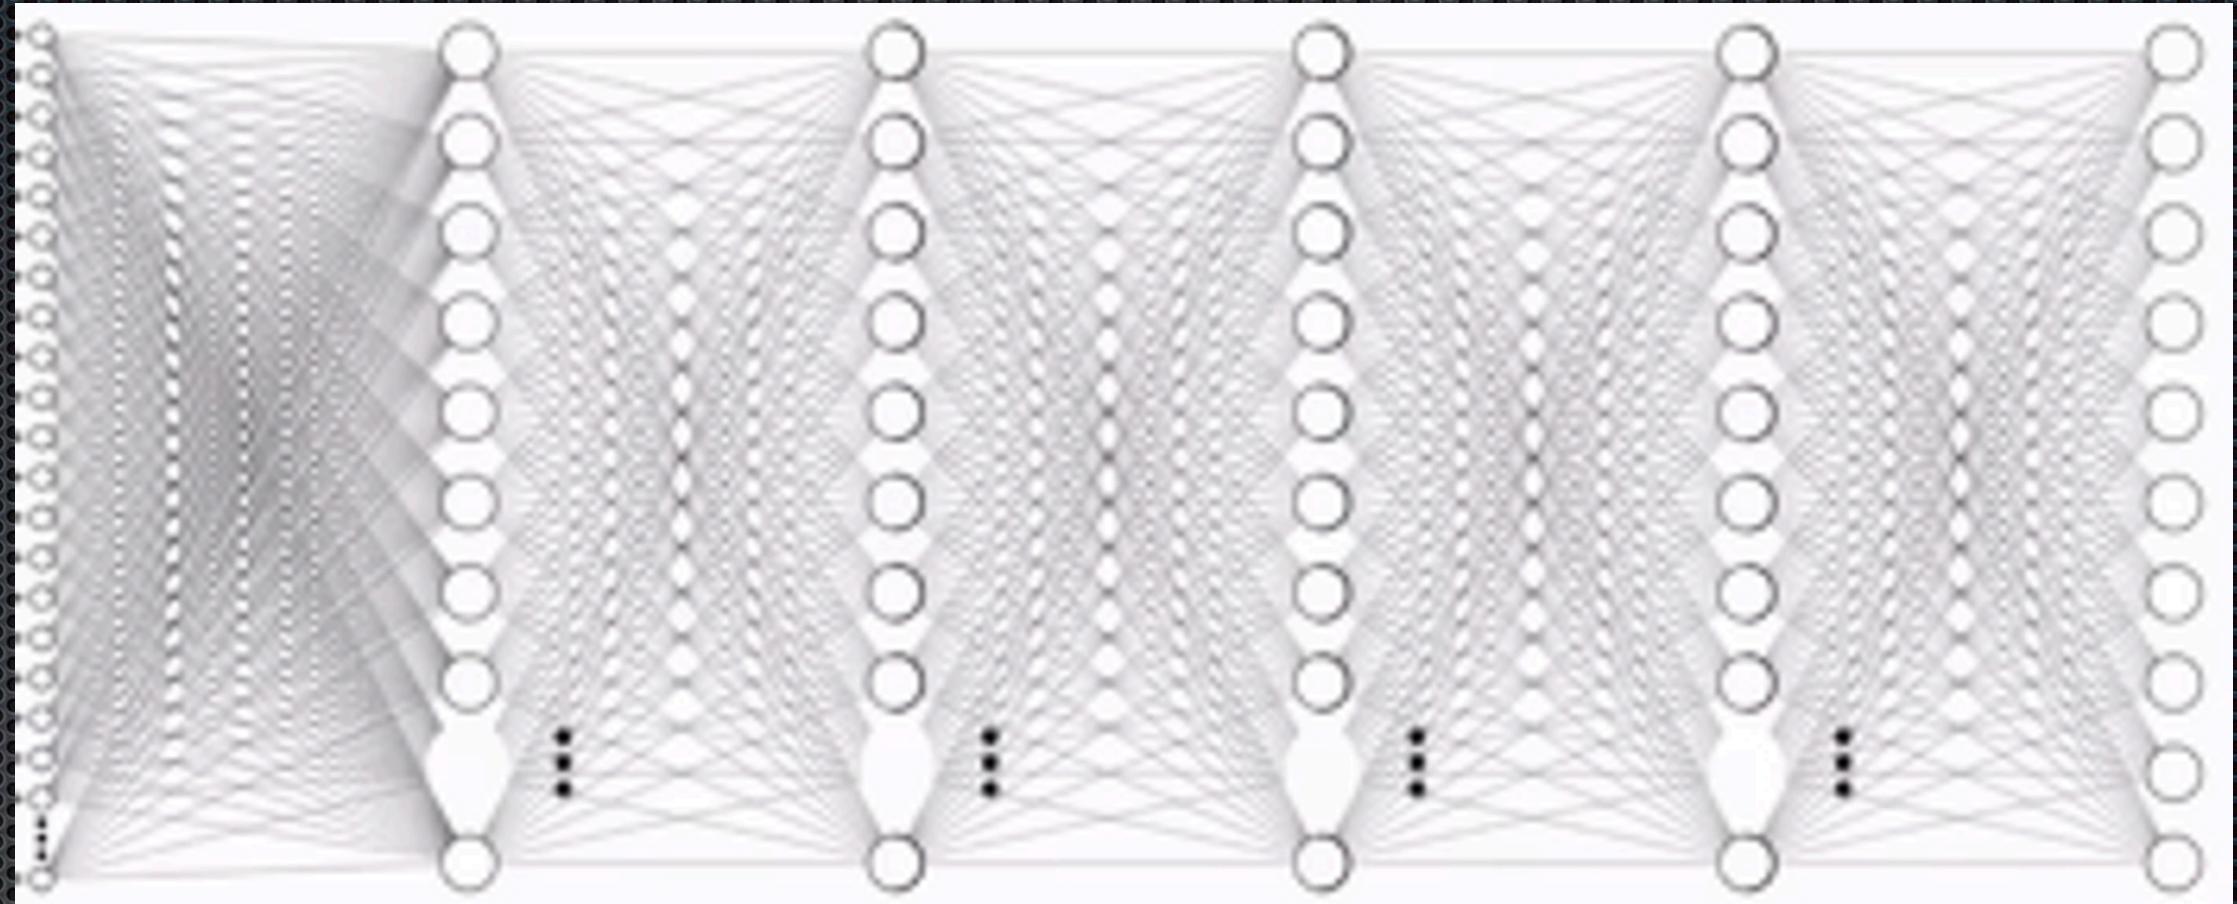
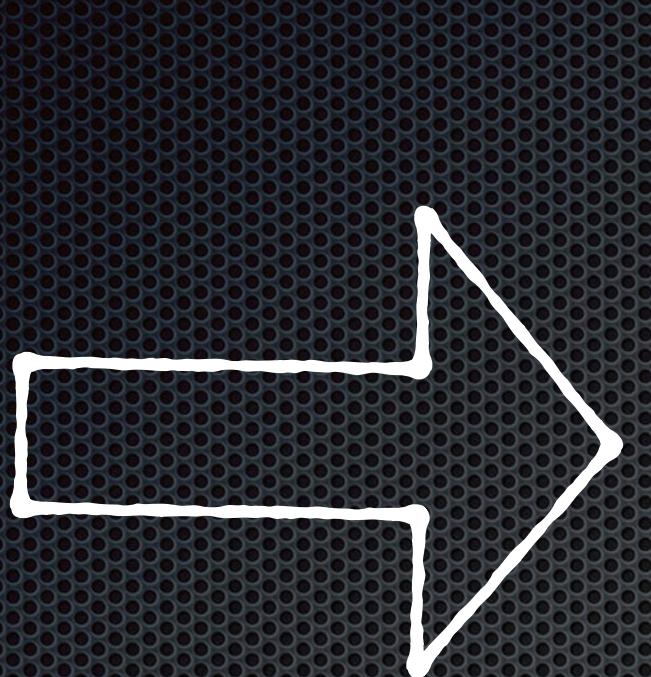
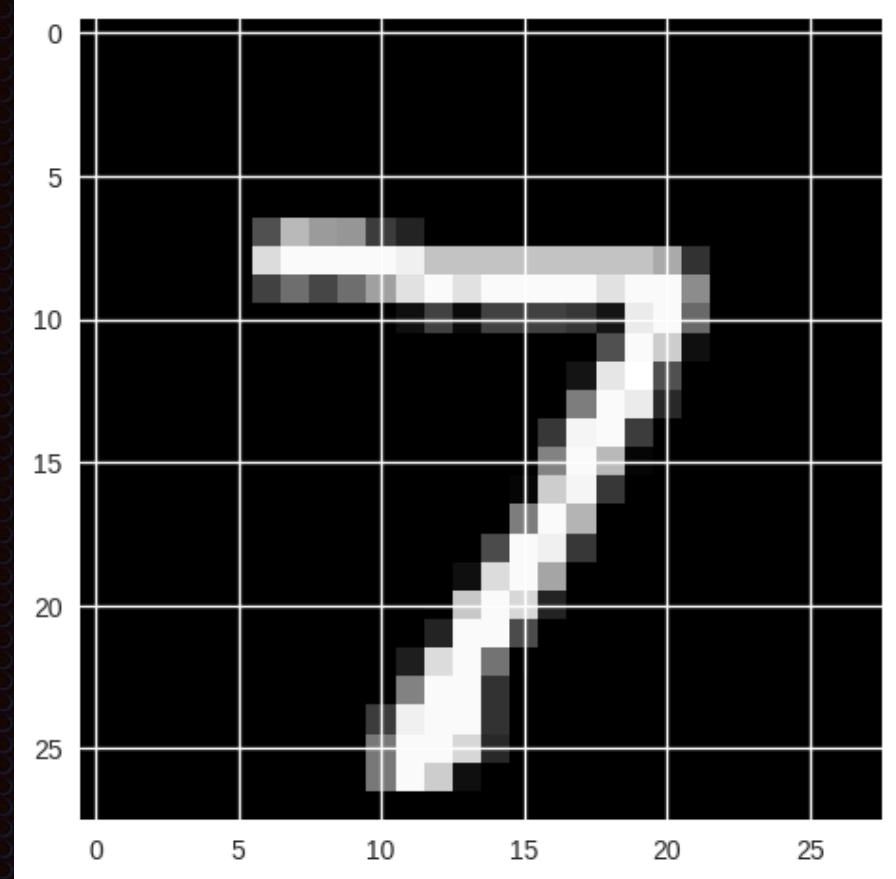
Introduction

Why do we need CNN models ? DNNs can also analyze images !!

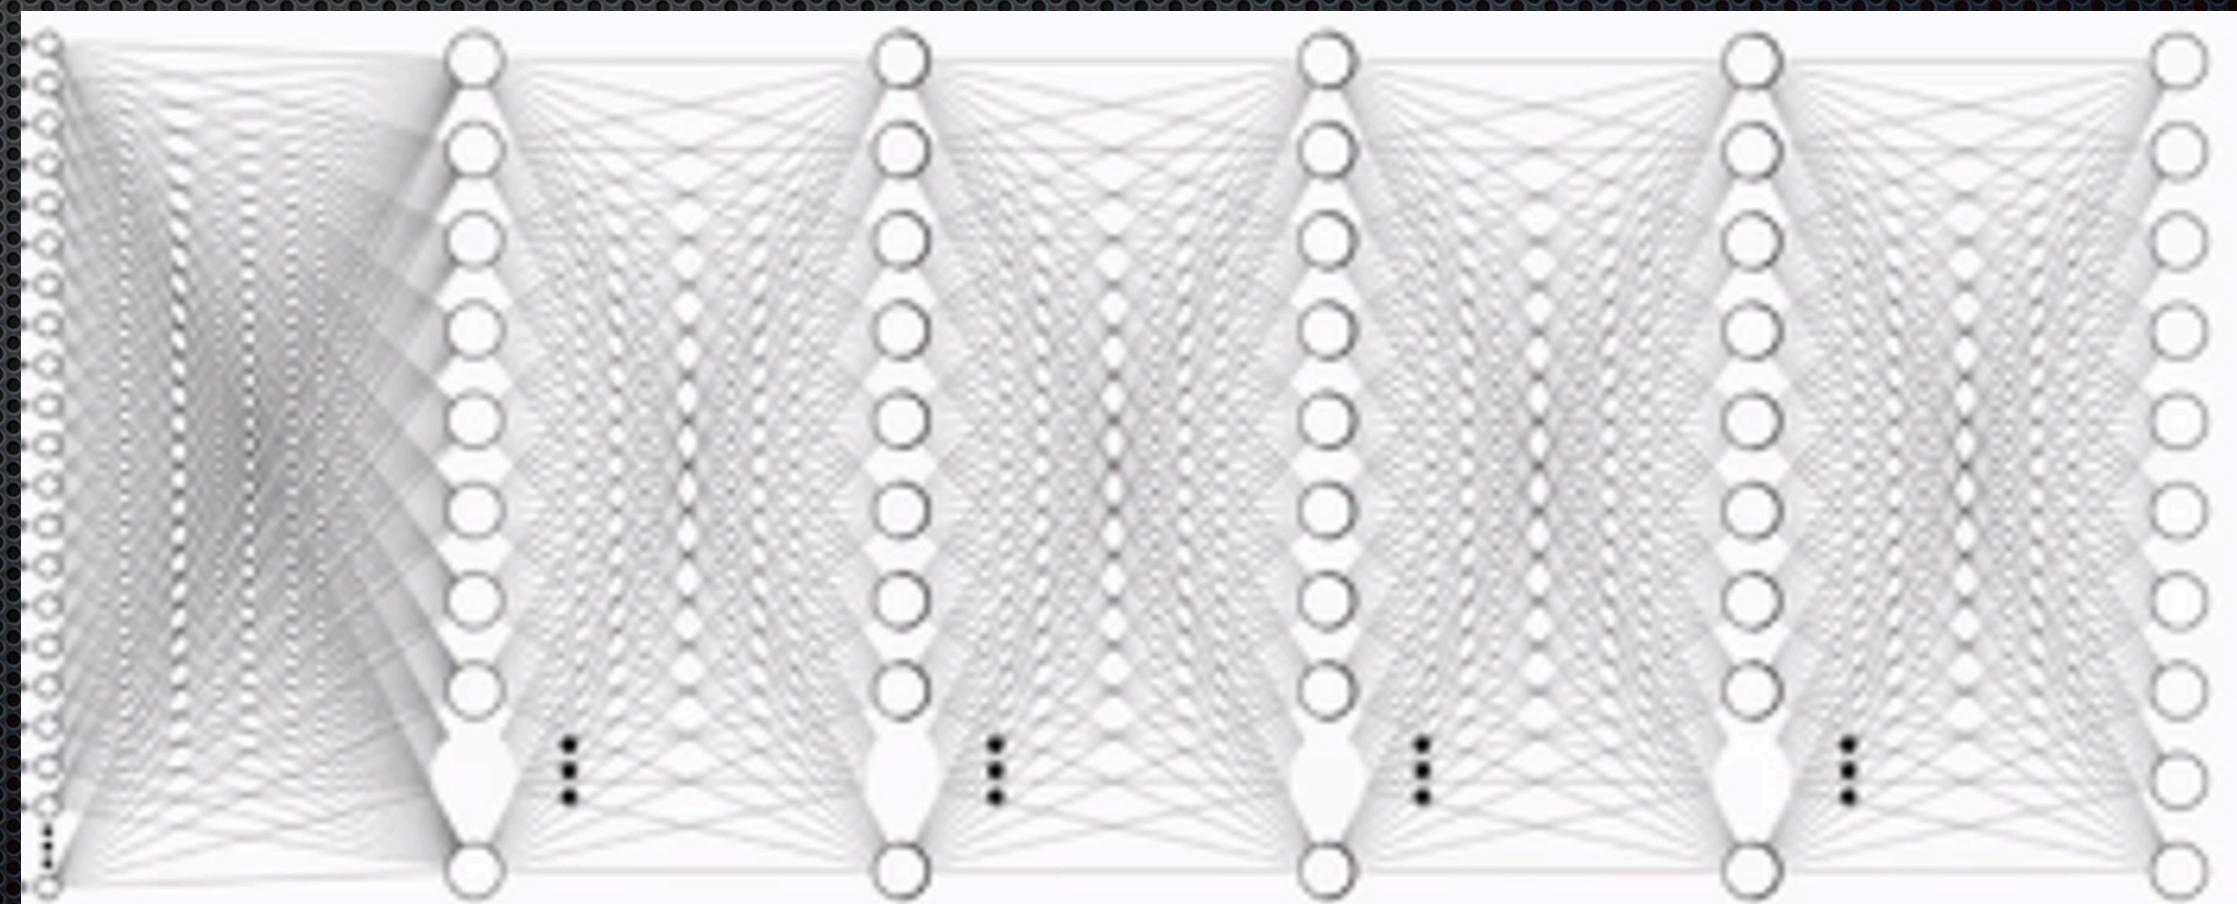
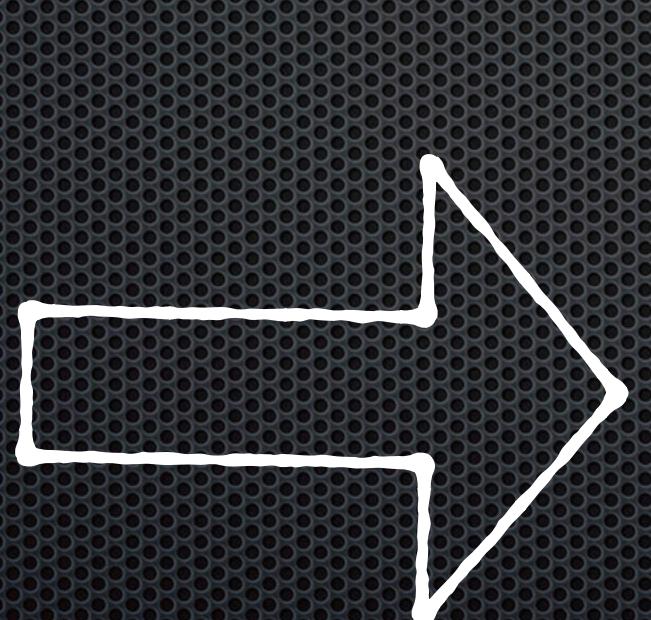
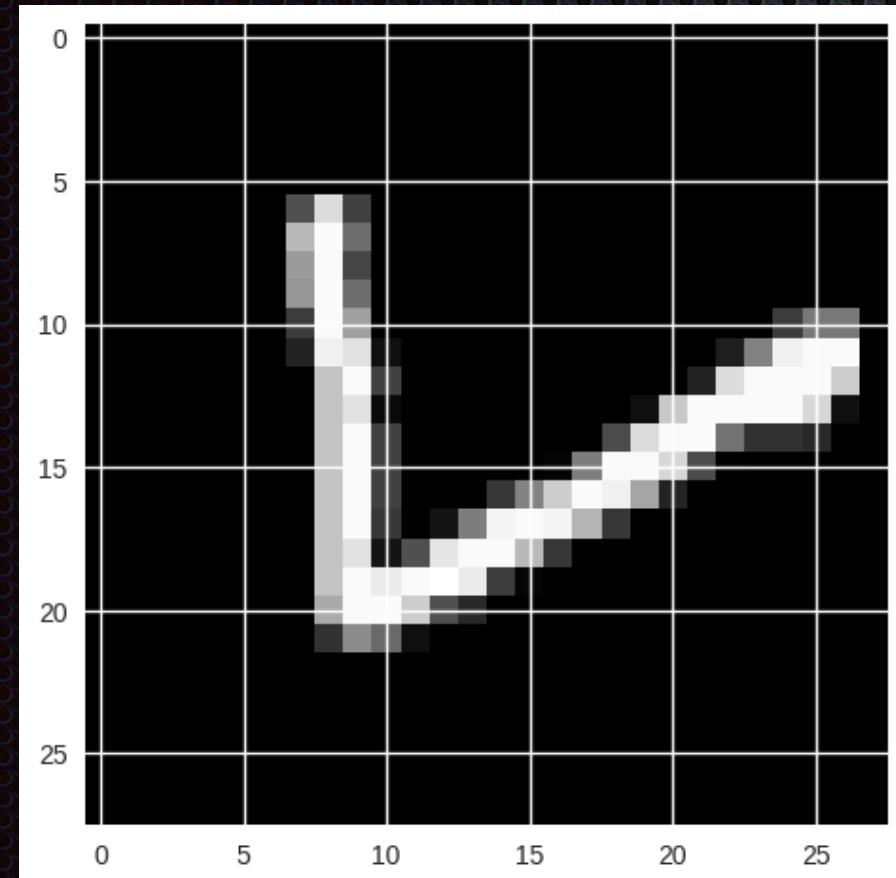
Introduction

Trained DNN from lecture 3

Why do we need CNN models ? DNNs can also analyze images !!



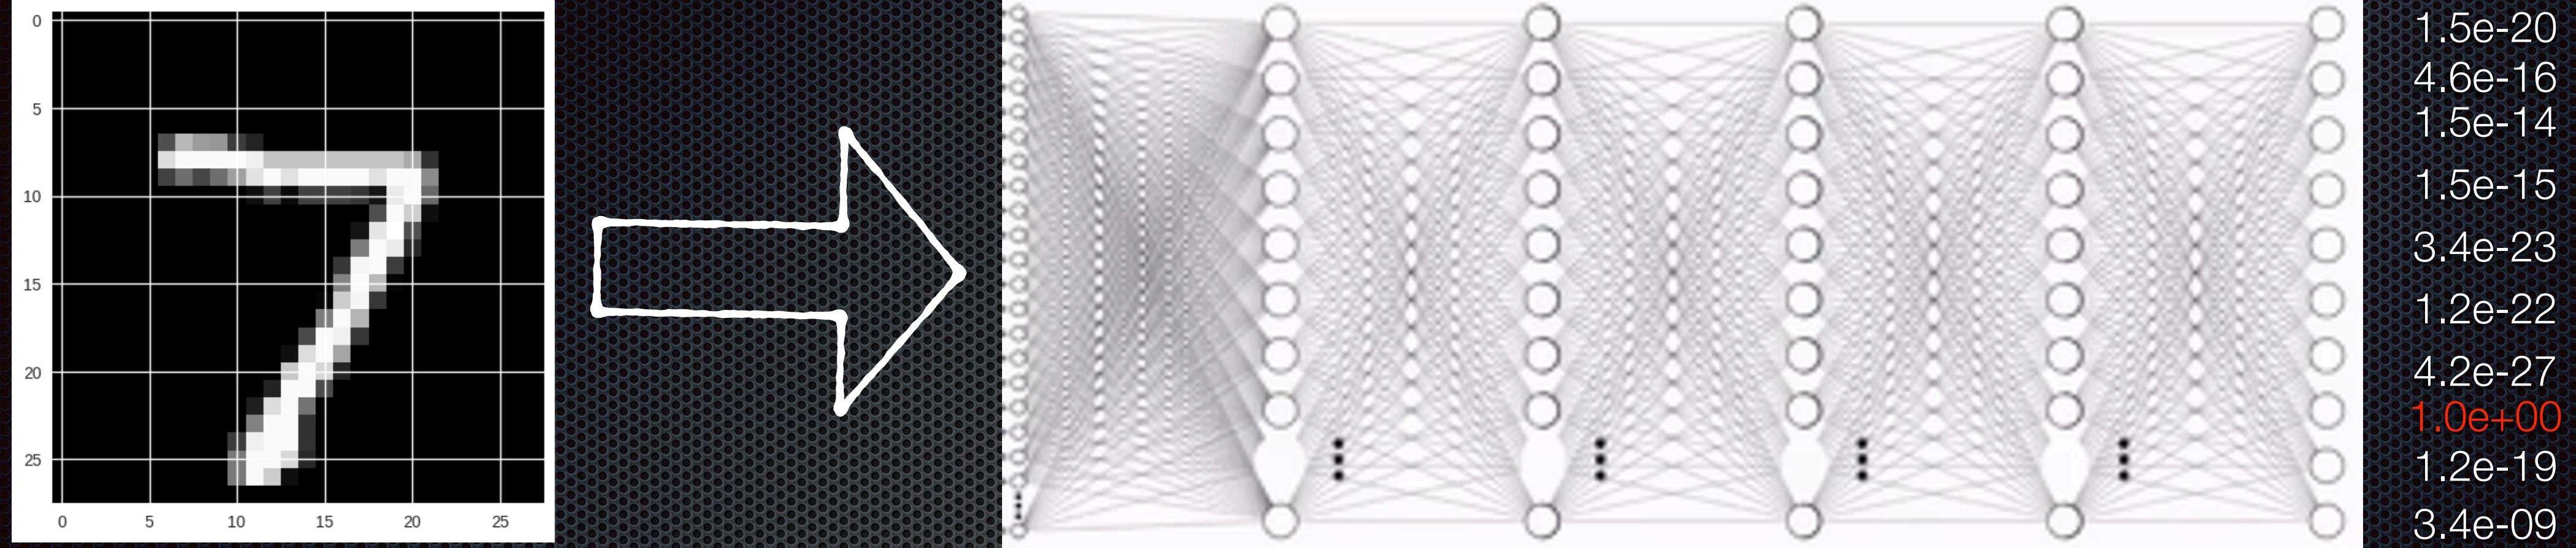
$1.5e-20$
 $4.6e-16$
 $1.5e-14$
 $1.5e-15$
 $3.4e-23$
 $1.2e-22$
 $4.2e-27$
 $1.0e+00$
 $1.2e-19$
 $3.4e-09$



$2.3e-04$
 $1.8e-06$
 $2.1e-04$
 $4.3e-07$
 $9.5e-06$
 $1.0e-02$
 $9.7e-01$
 $1.2e-07$
 $1.8e-02$
 $1.3e-06$

Introduction

Why do we need CNN models ? DNNs can also analyze images !!

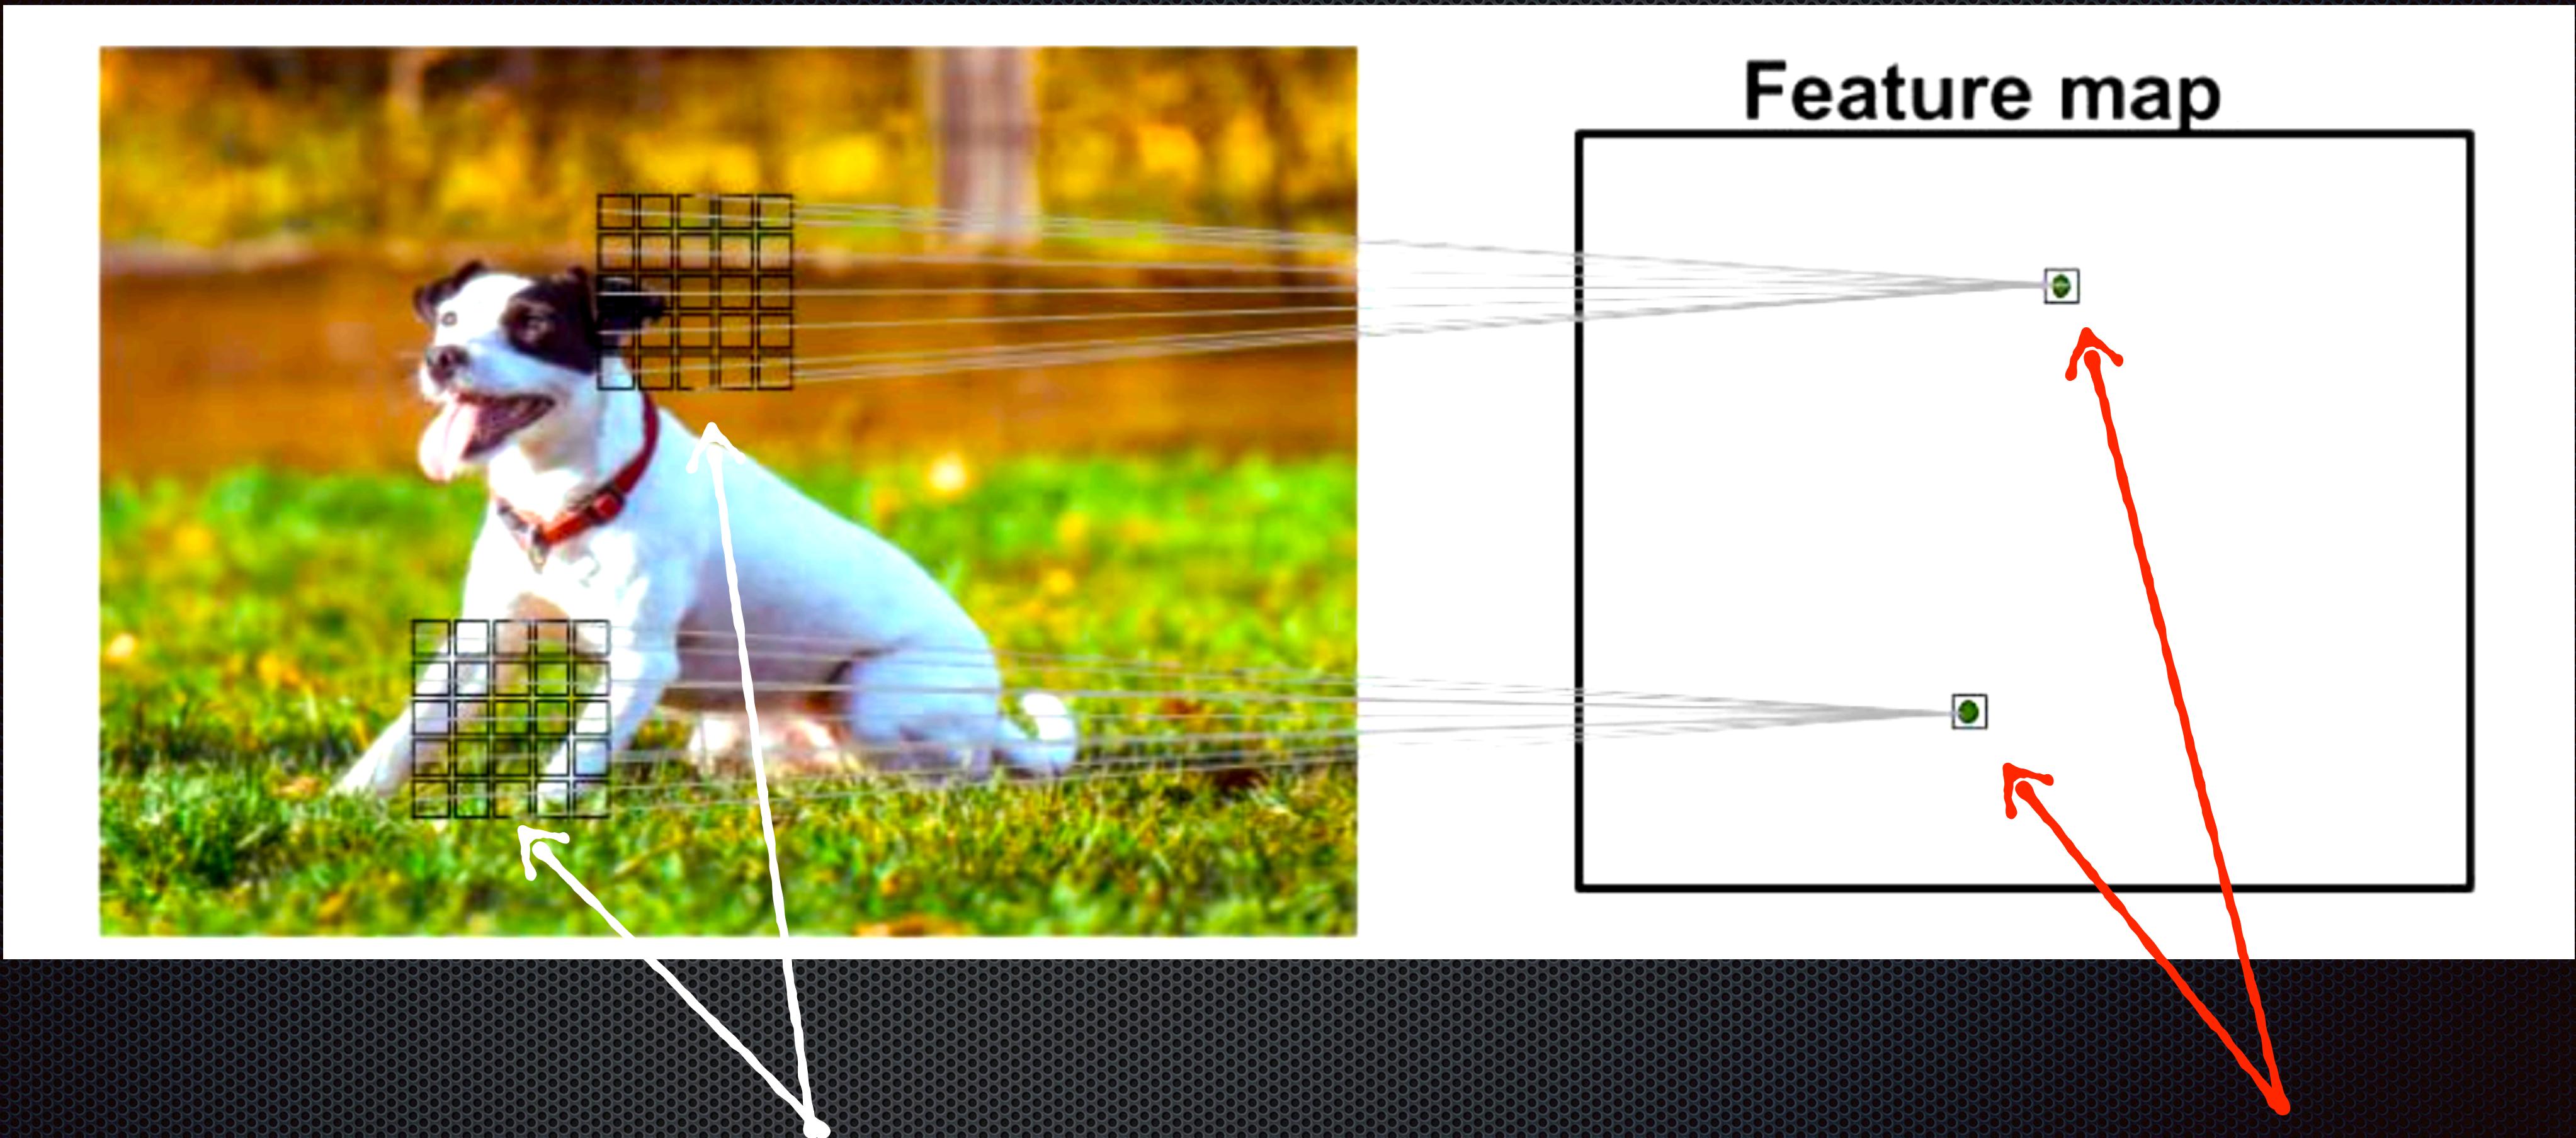


As a fully connected model

$$\text{Neuron Output} = \text{Activation function} \left(\sum_m \omega_m x_m + b \right)$$

The model learns the correlation between all pixels into the image and construct a global information about the image

Introduction



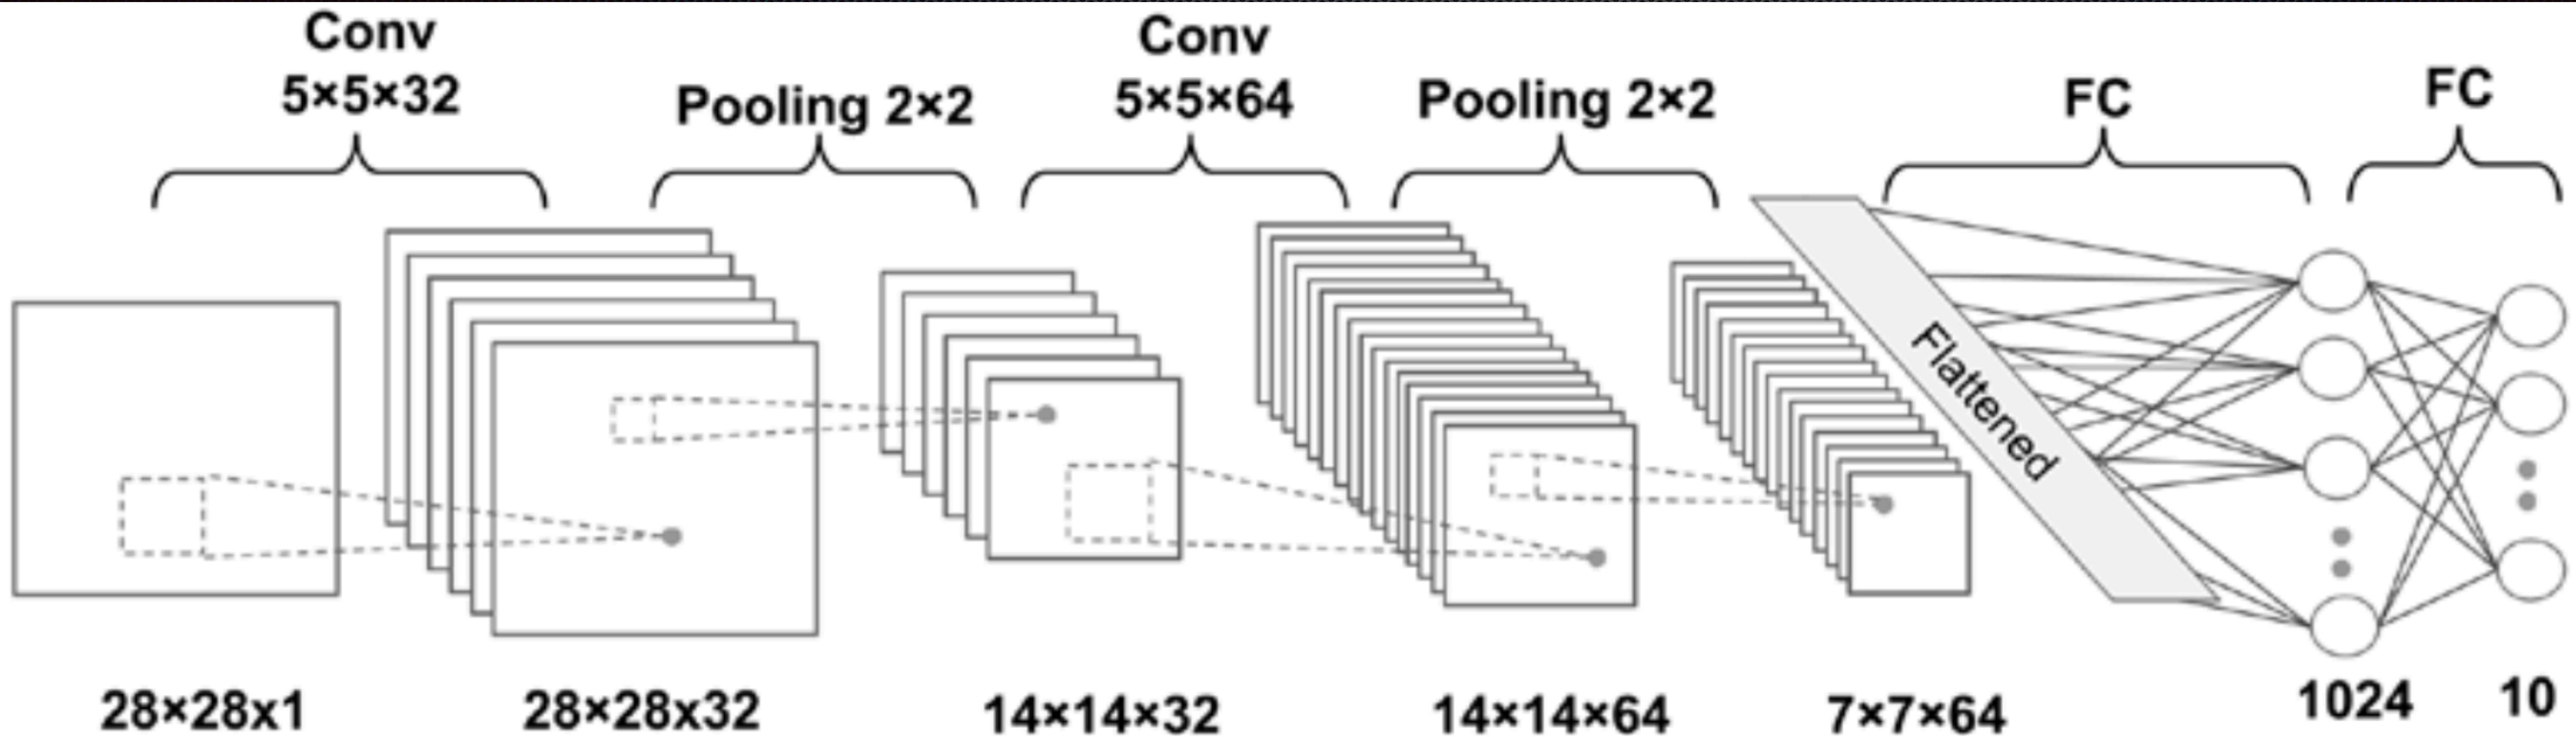
Filters to capture the **local features** of the images

Filters share their weights internally but not with other filters

Disconnected
Patches

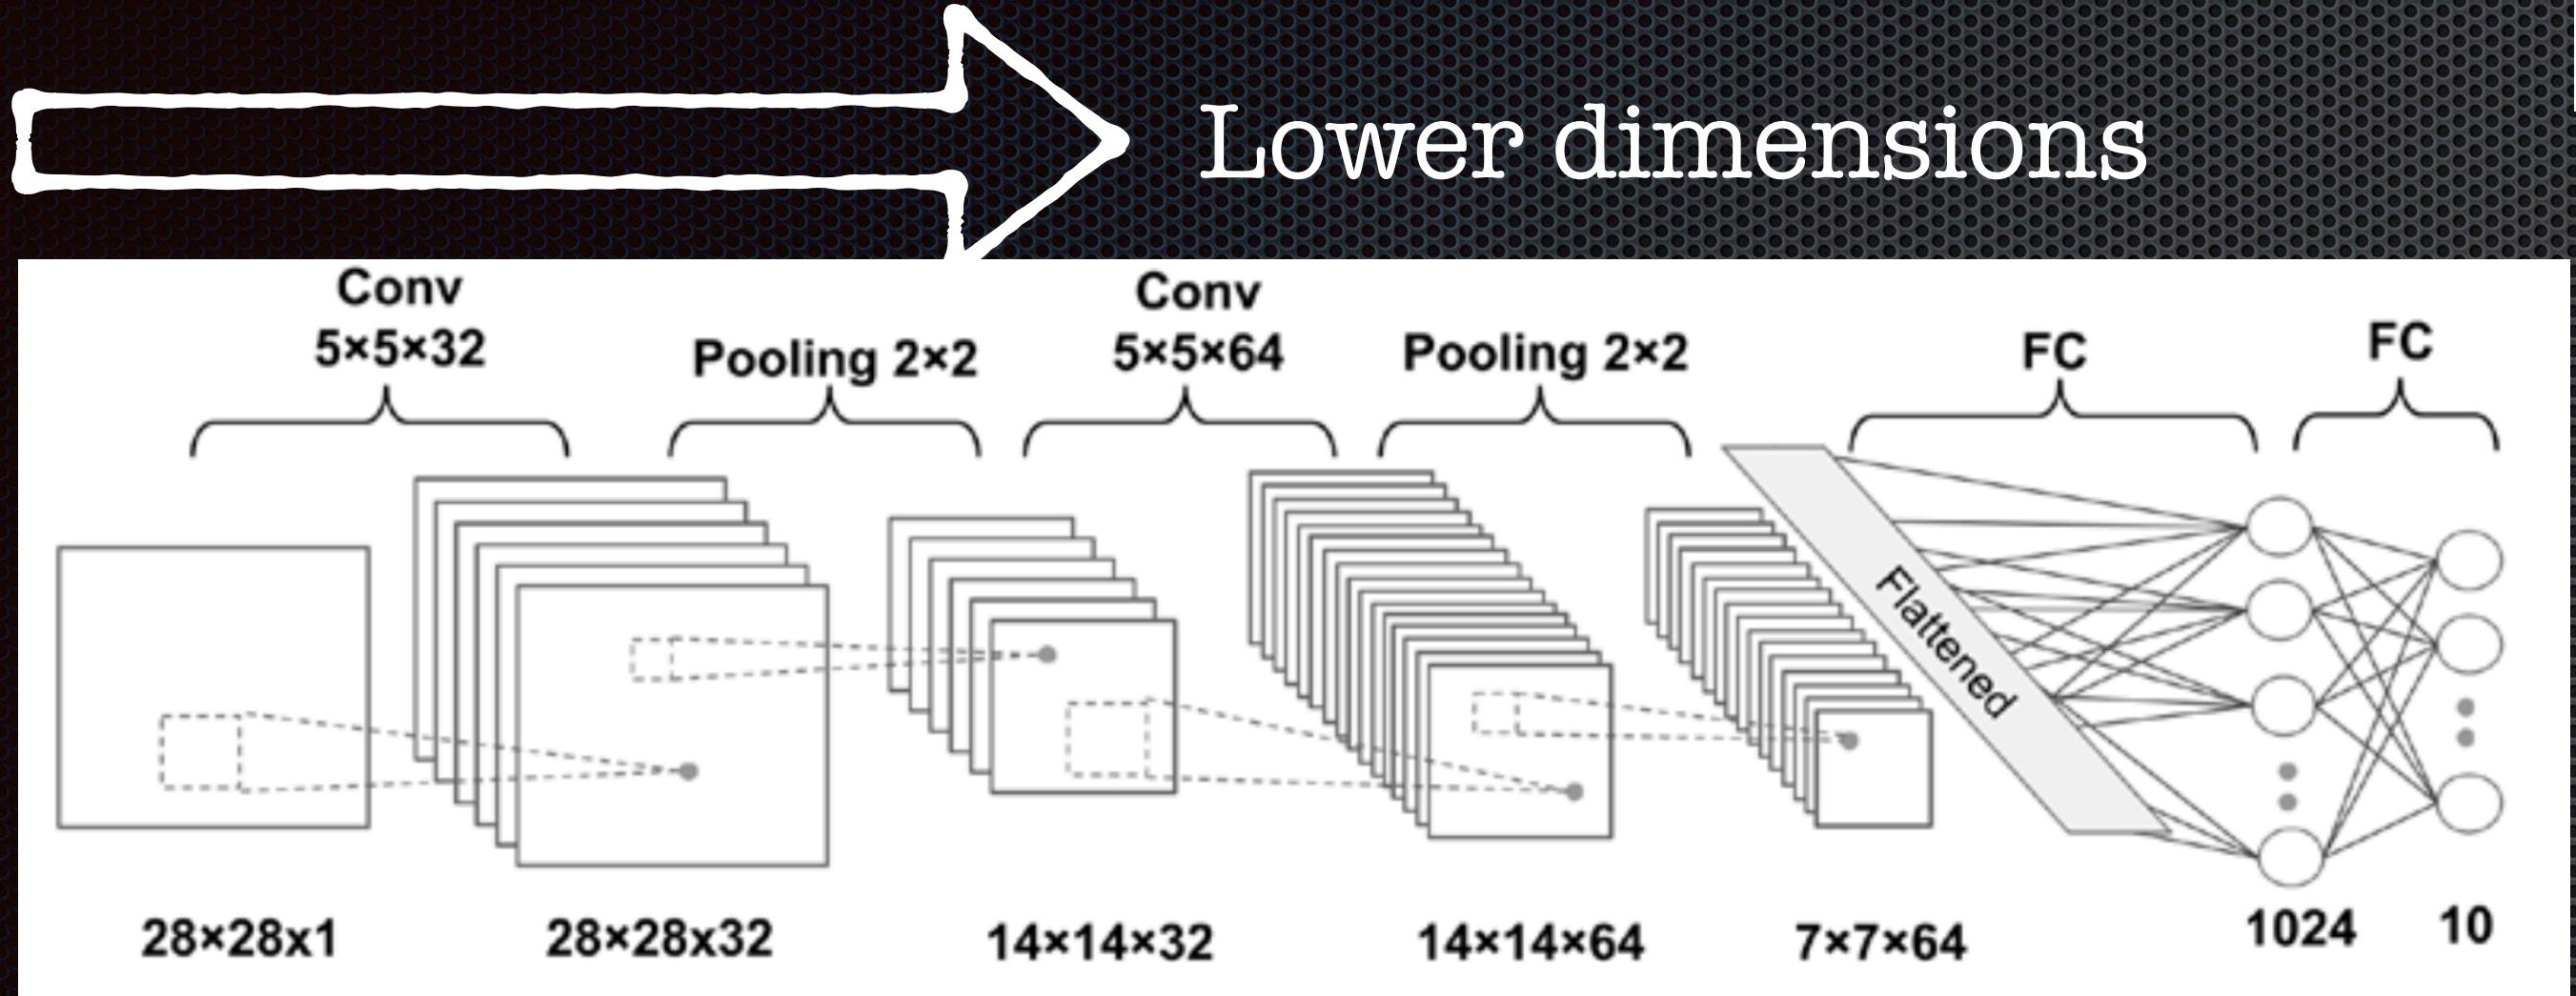
Introduction

Example of deep CNN model

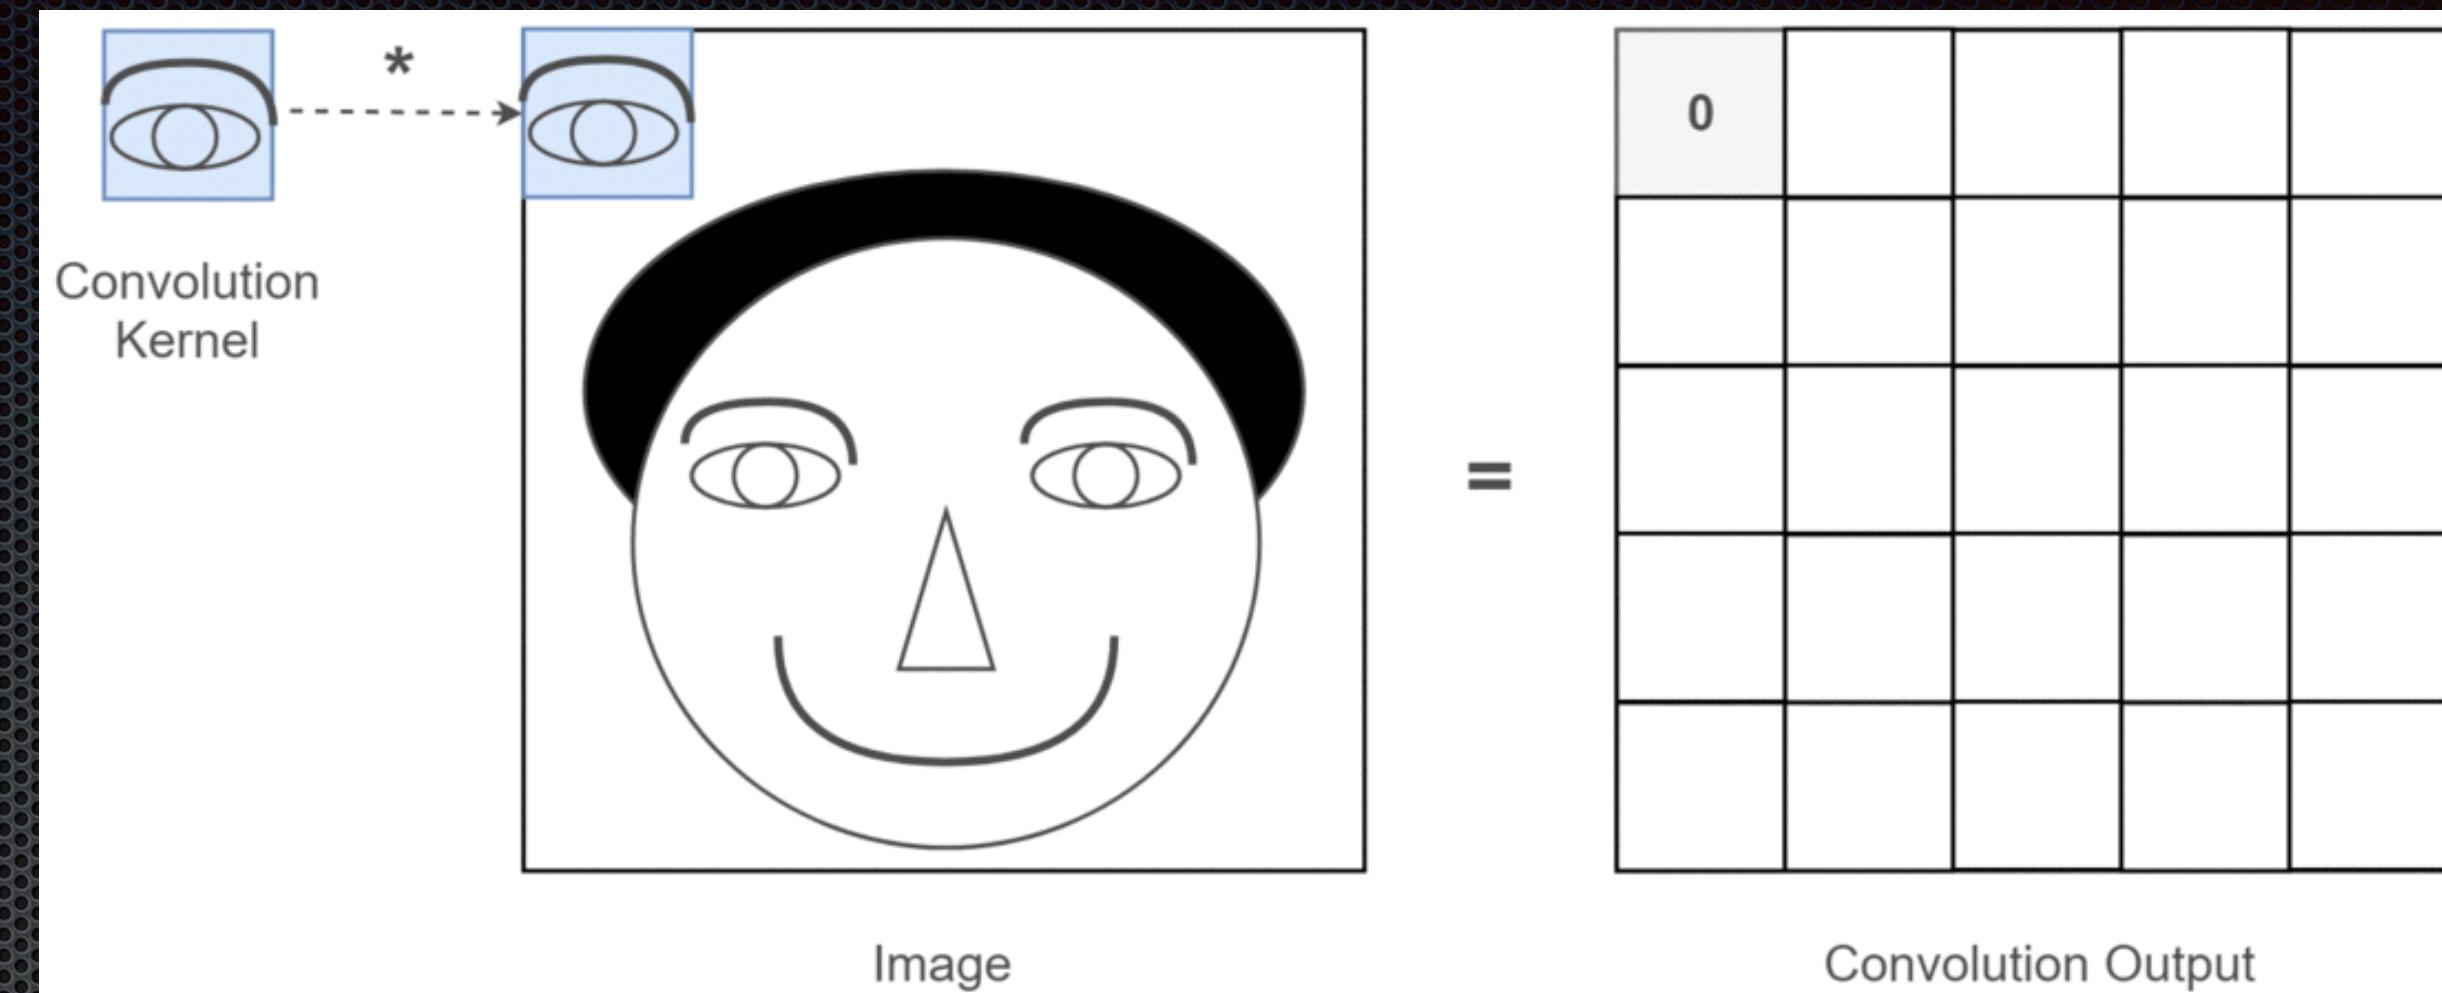


CNN = Convolution layers + Fully connected layers

All together



Lower dimensions



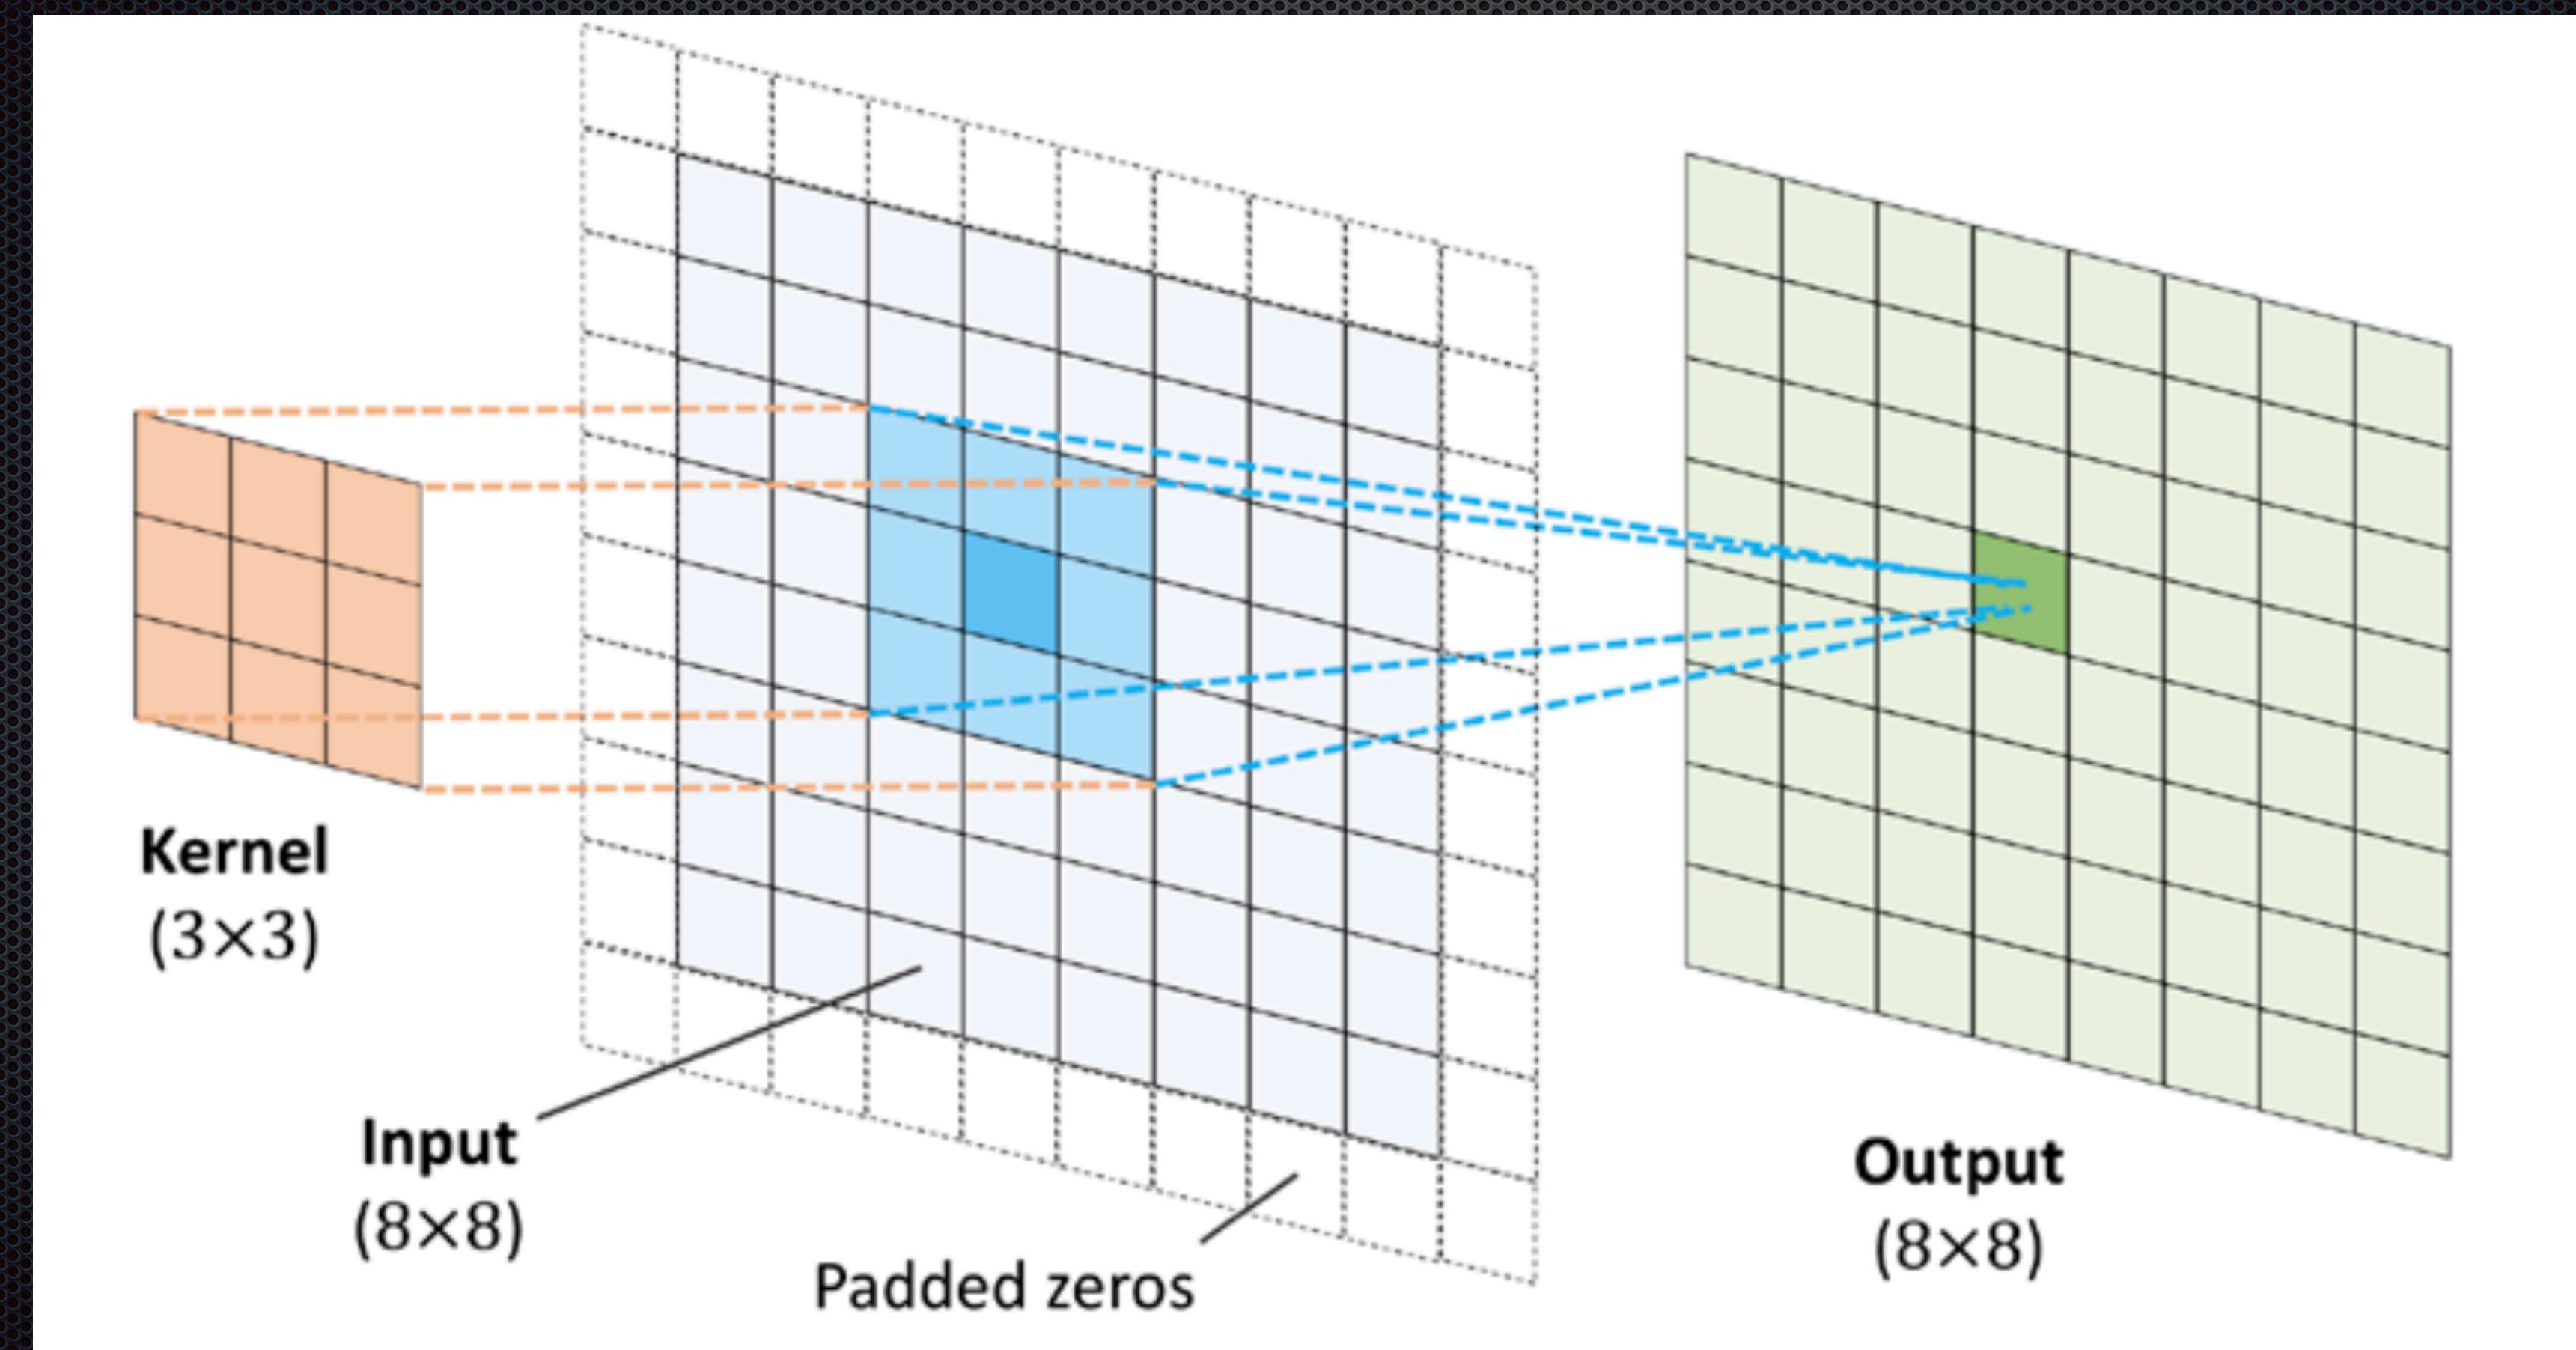
Using many filters to learn more local features, e.g. nose, eyes, etc

After the captured local information is mapped onto the decomposed latent space we use a fully connected layers to analyze these information and learn a global information about the input image

See lecture 4

Convolution in two dimensions

$$y[i, j] = \sum_{k_1=0}^{k_1=m_1-1} \sum_{k_2=0}^{k_2=m_2-1} x_p[i + m_1 - k_1, j + m_2 - k_2] \omega[k_1, k_2]$$



Convolution in two dimensions

0 ₂	0 ₀	0 ₁	0	0	0	0
0 ₁	2 ₀	2 ₀	3	3	3	0
0 ₀	0 ₁	1 ₁	3	0	3	0
0	2	3	0	1	3	0
0	3	3	2	1	2	0
0	3	3	0	2	3	0
0	0	0	0	0	0	0

1	6	5
7	10	9
7	10	8

Local weights sharing

Convolution output =

Size of input vector + 2*Padding - filter size

Stride

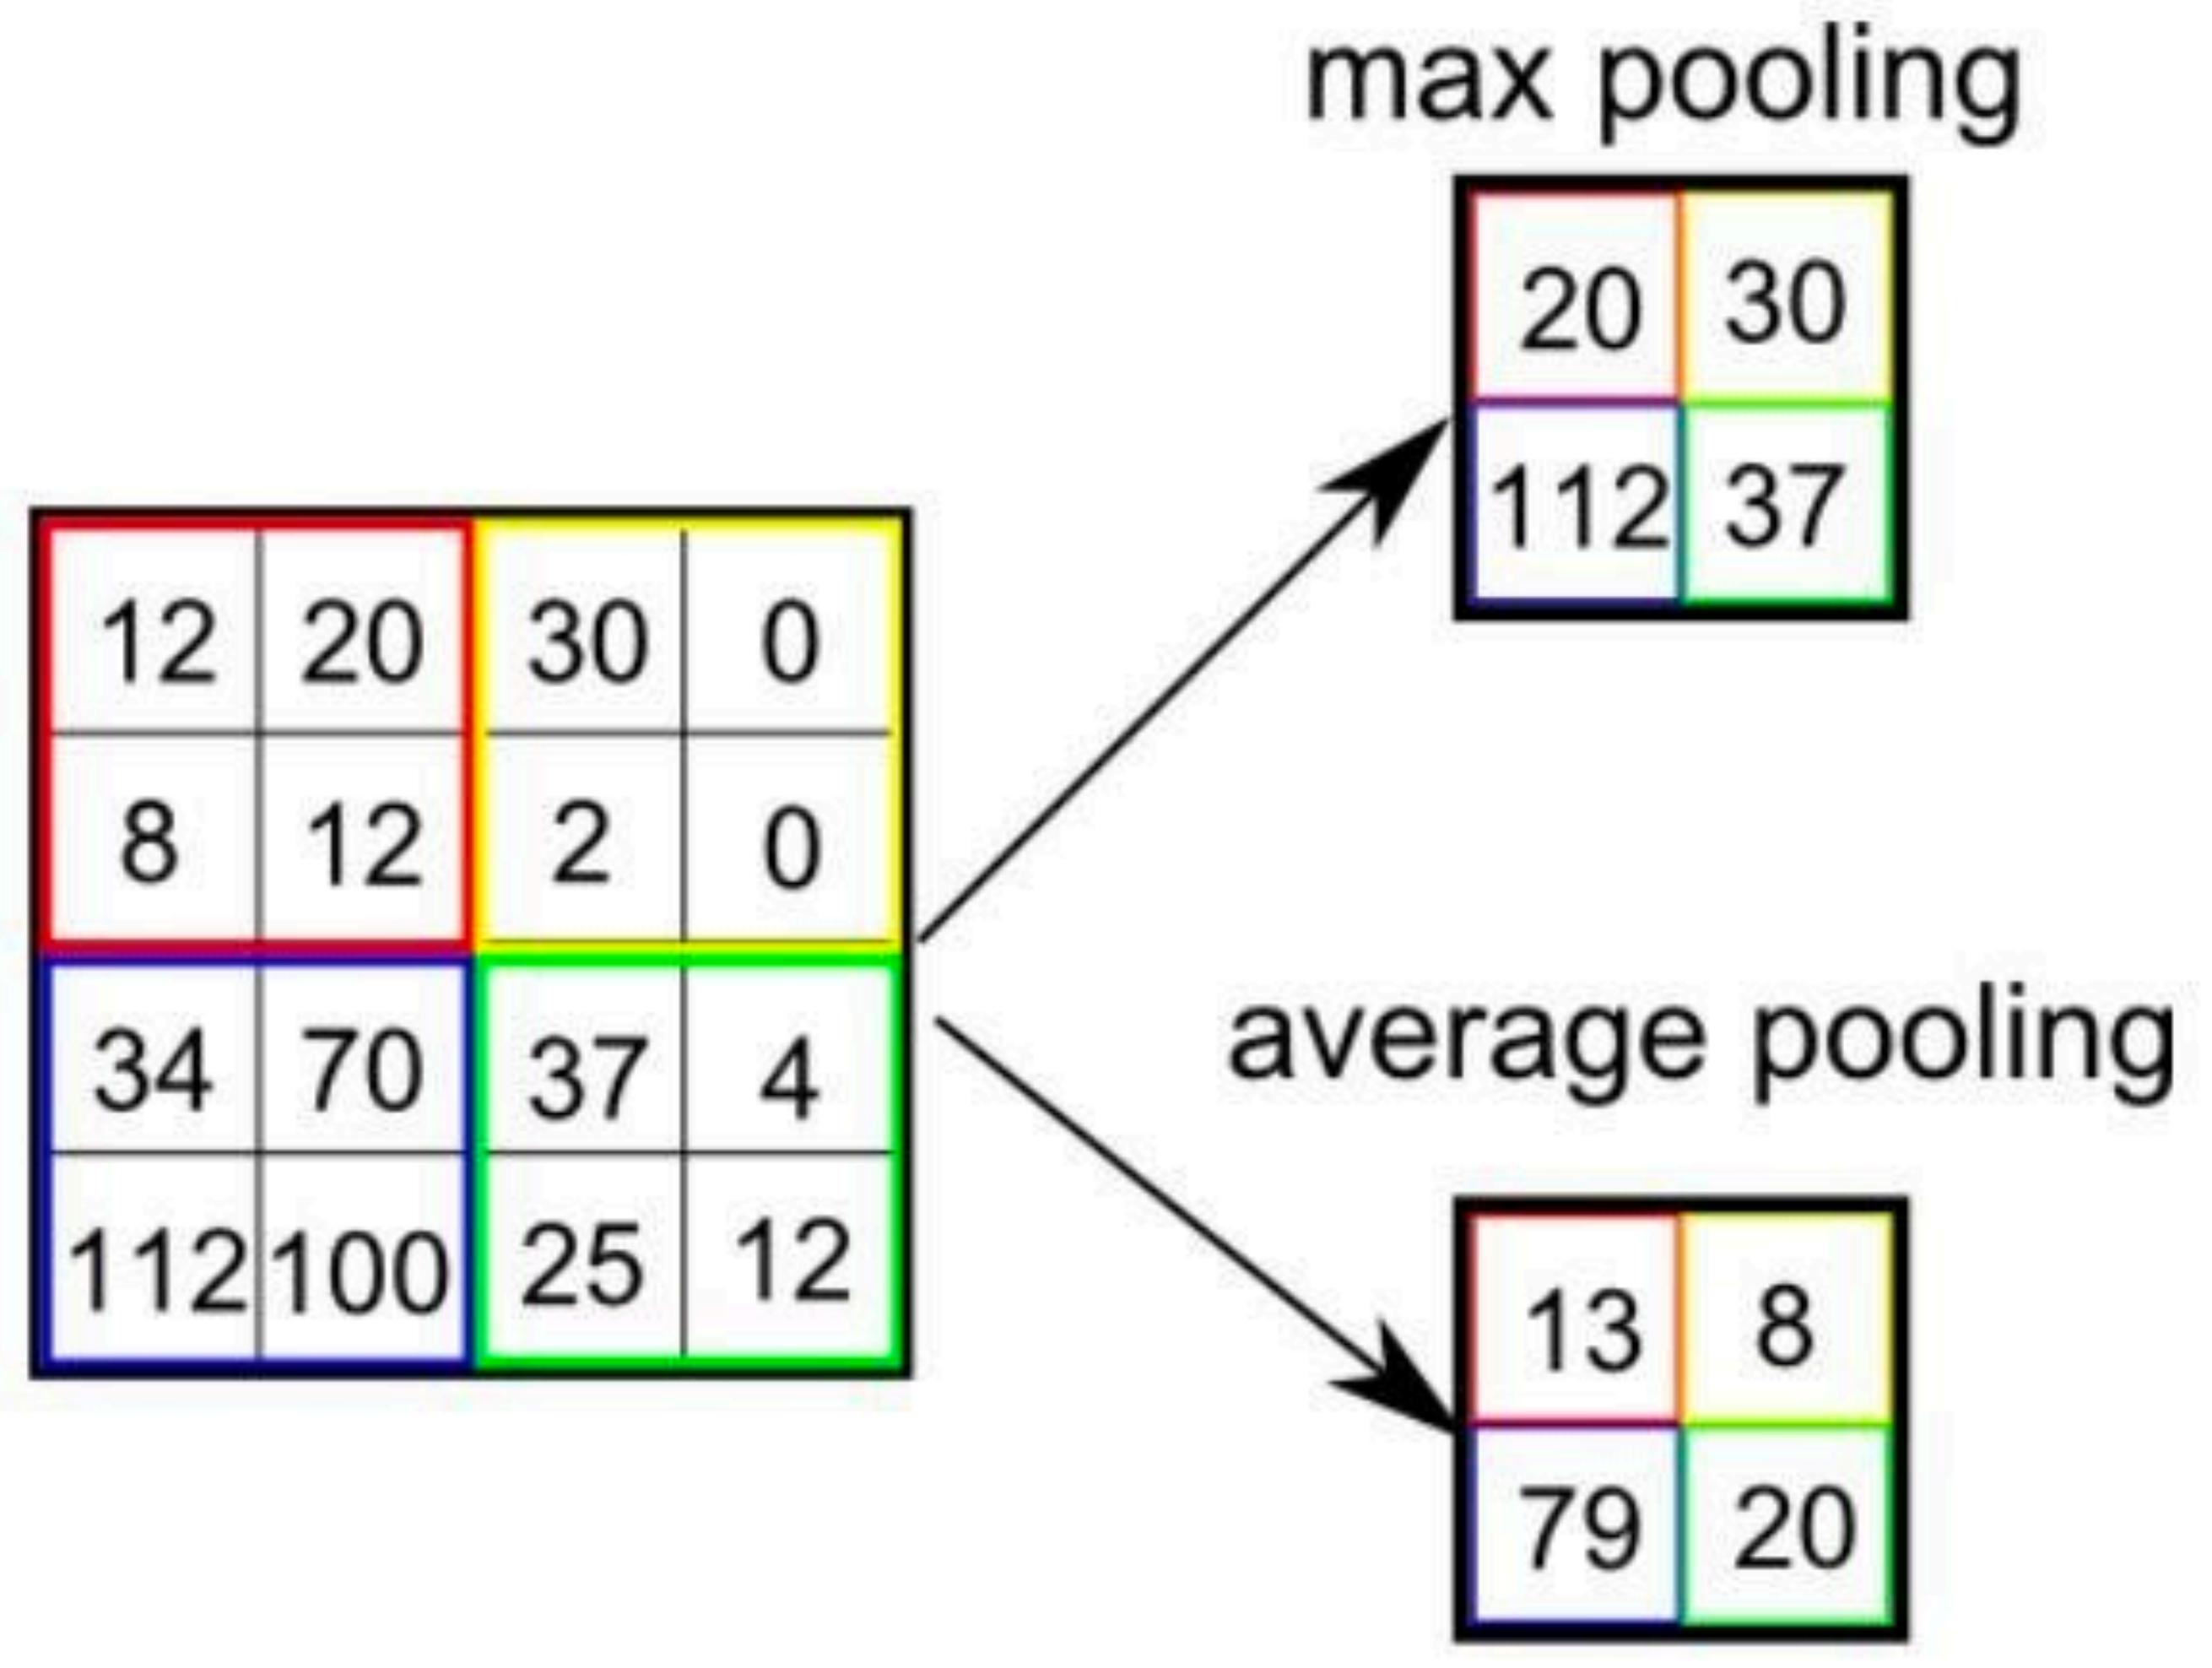
Convolution output =

$(5,5) + 2*(1,1) - (3,3)$
 $(2,2)$ + 1 = (3,3)

+1

Pooling

Pooling reduces the high dimensions Feature space to lower dimensions space (Latent space) to capture the characteristic information only

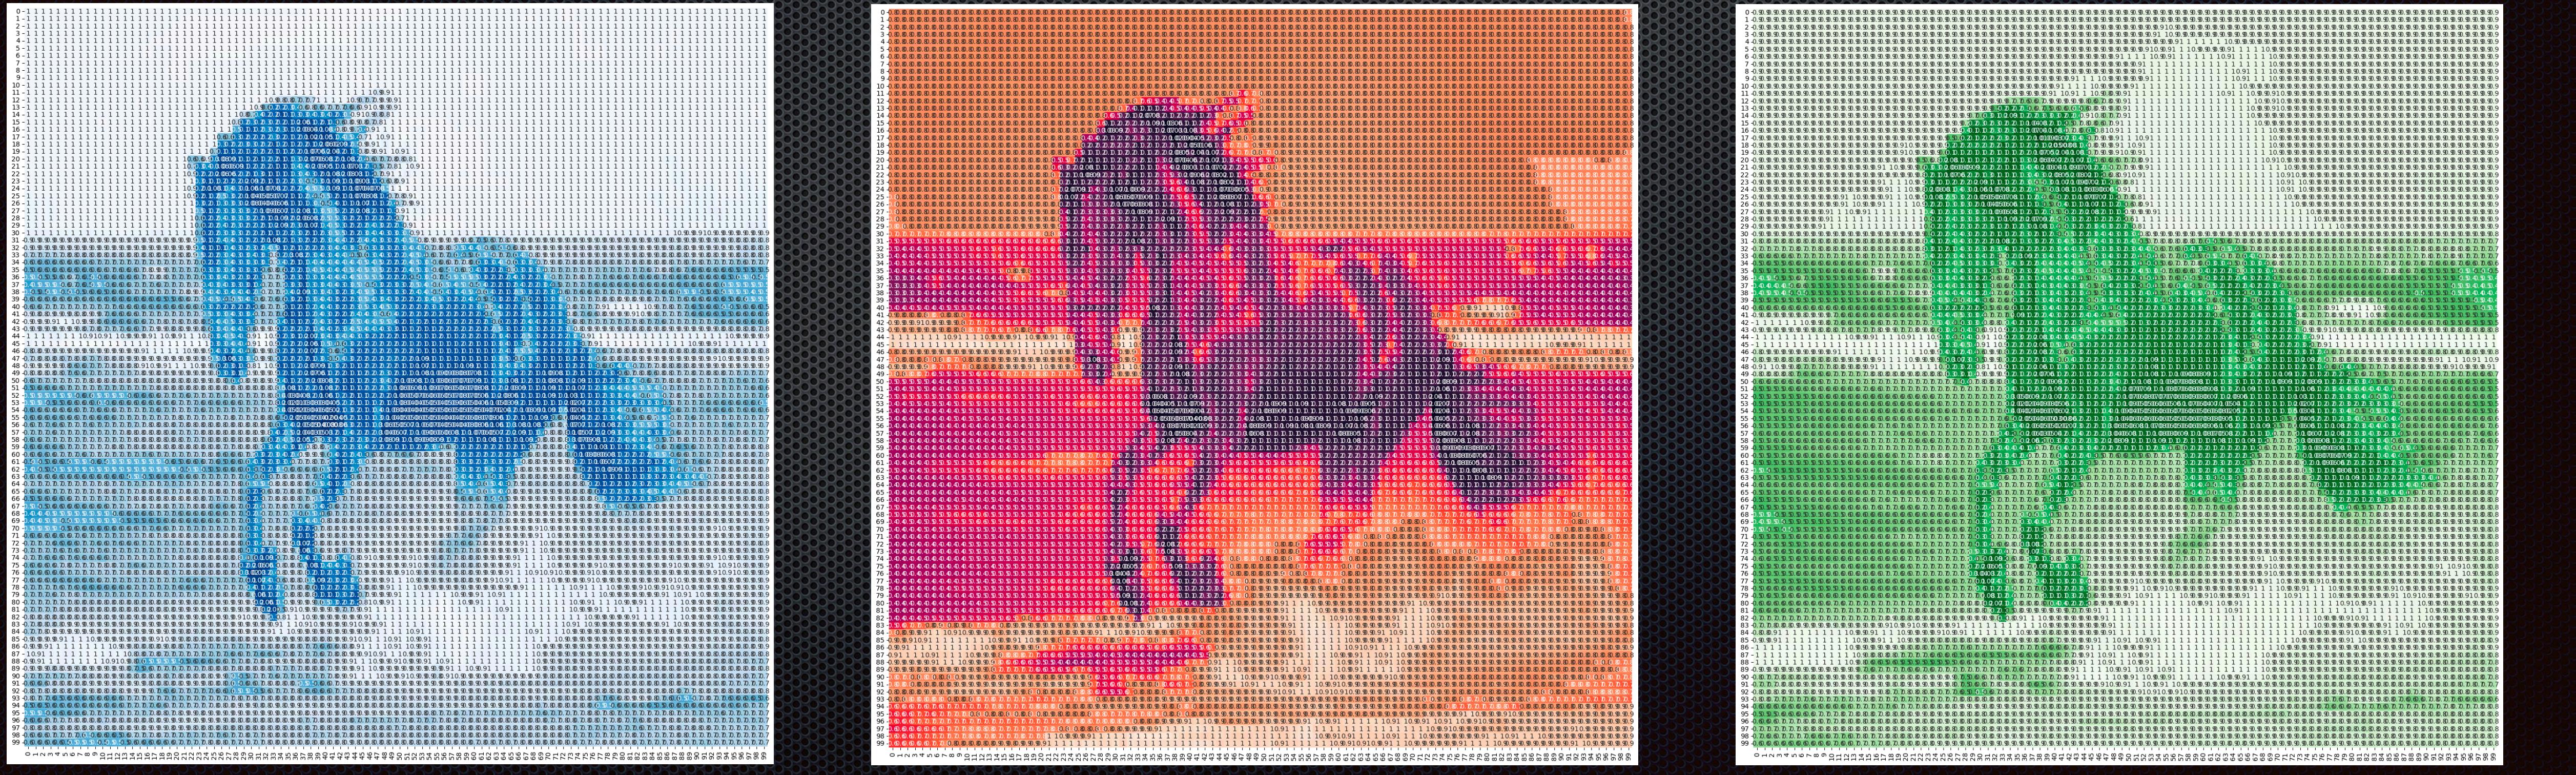
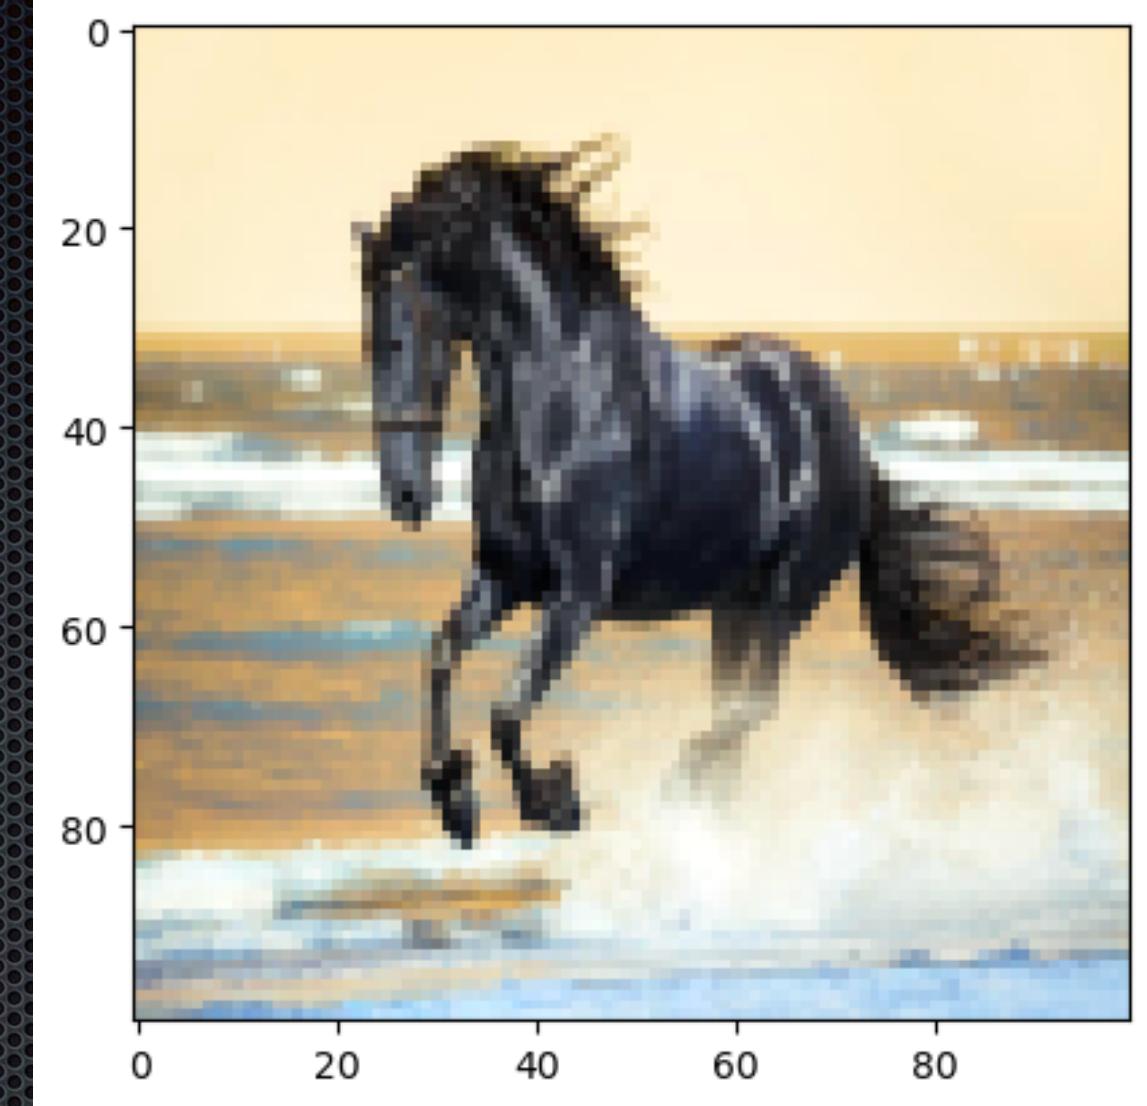


Dealing with colored images

Blue dimension
Dimension: (100,100,1)

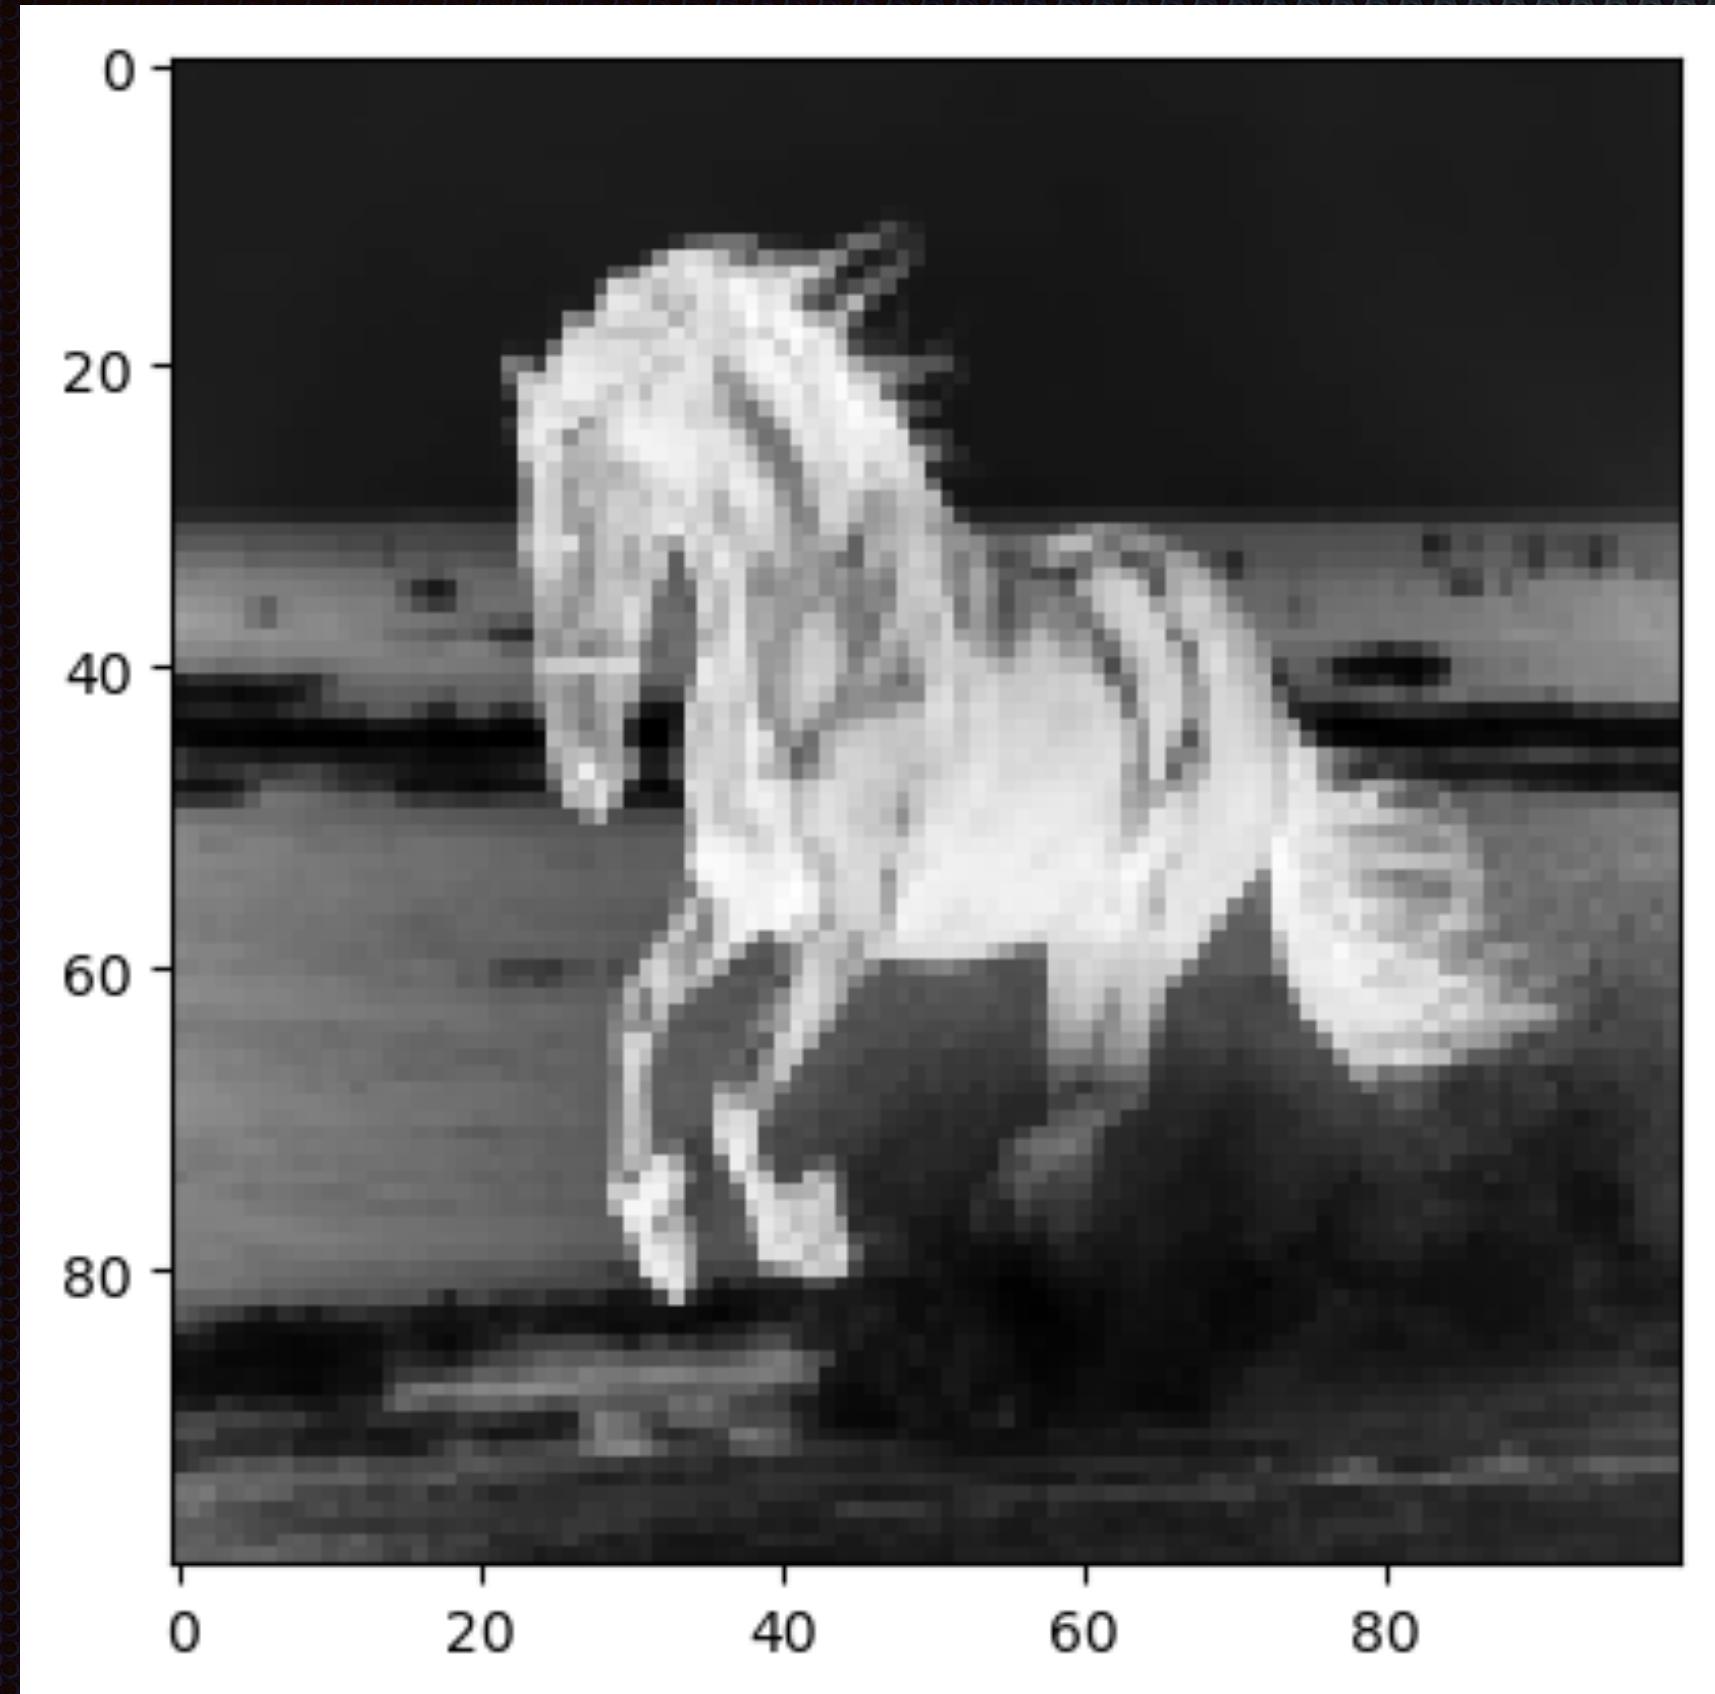
Red dimension
Dimension: (100,100,1)

Original image
Dimension: (100,100,3)



Dealing with colored images

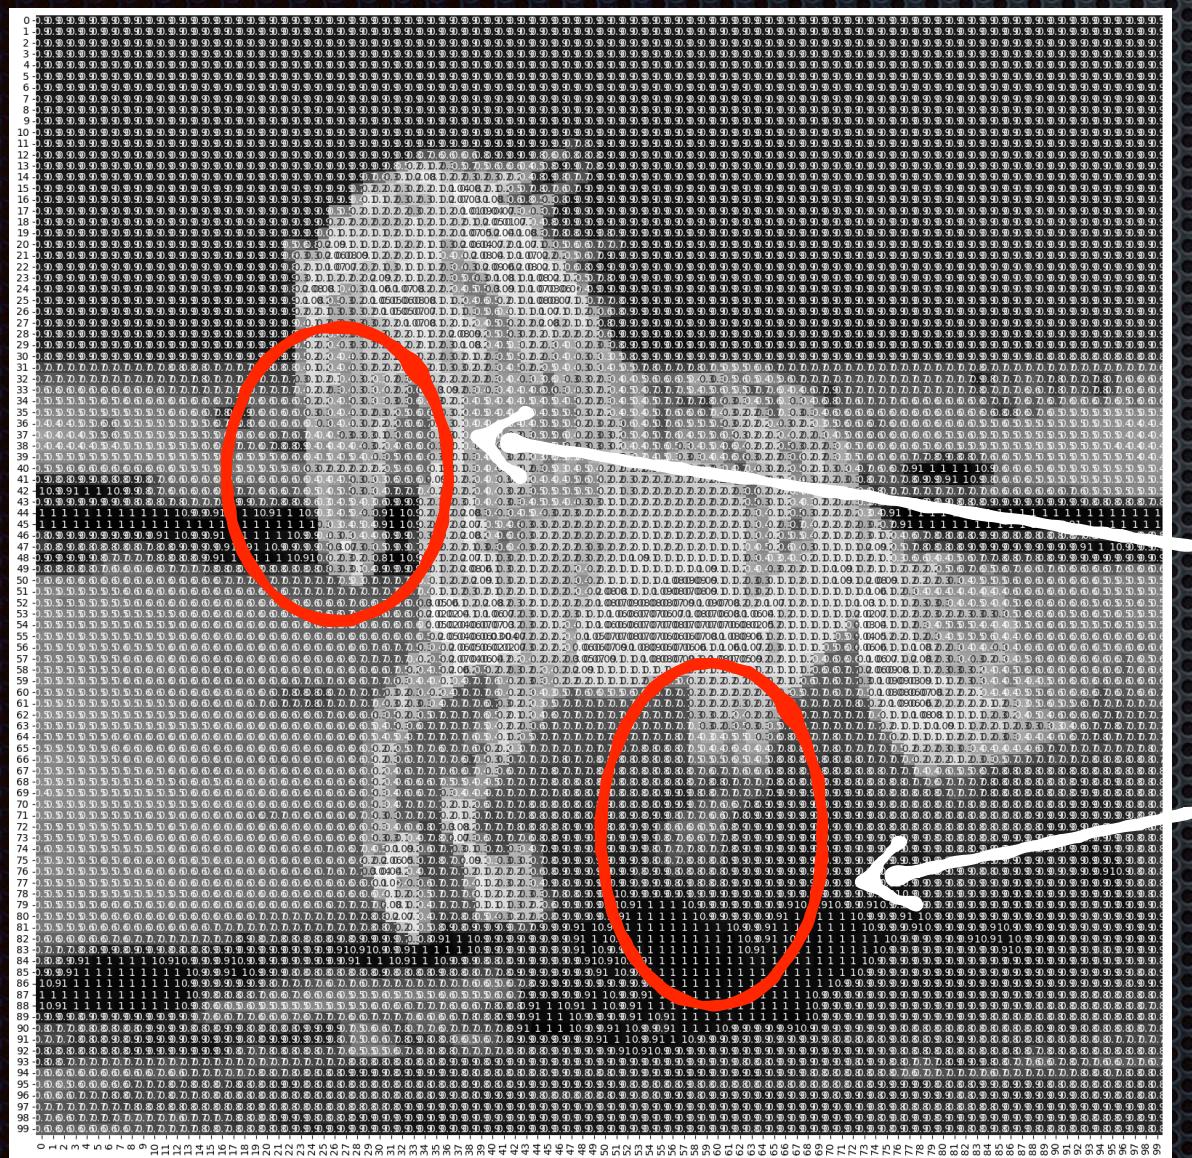
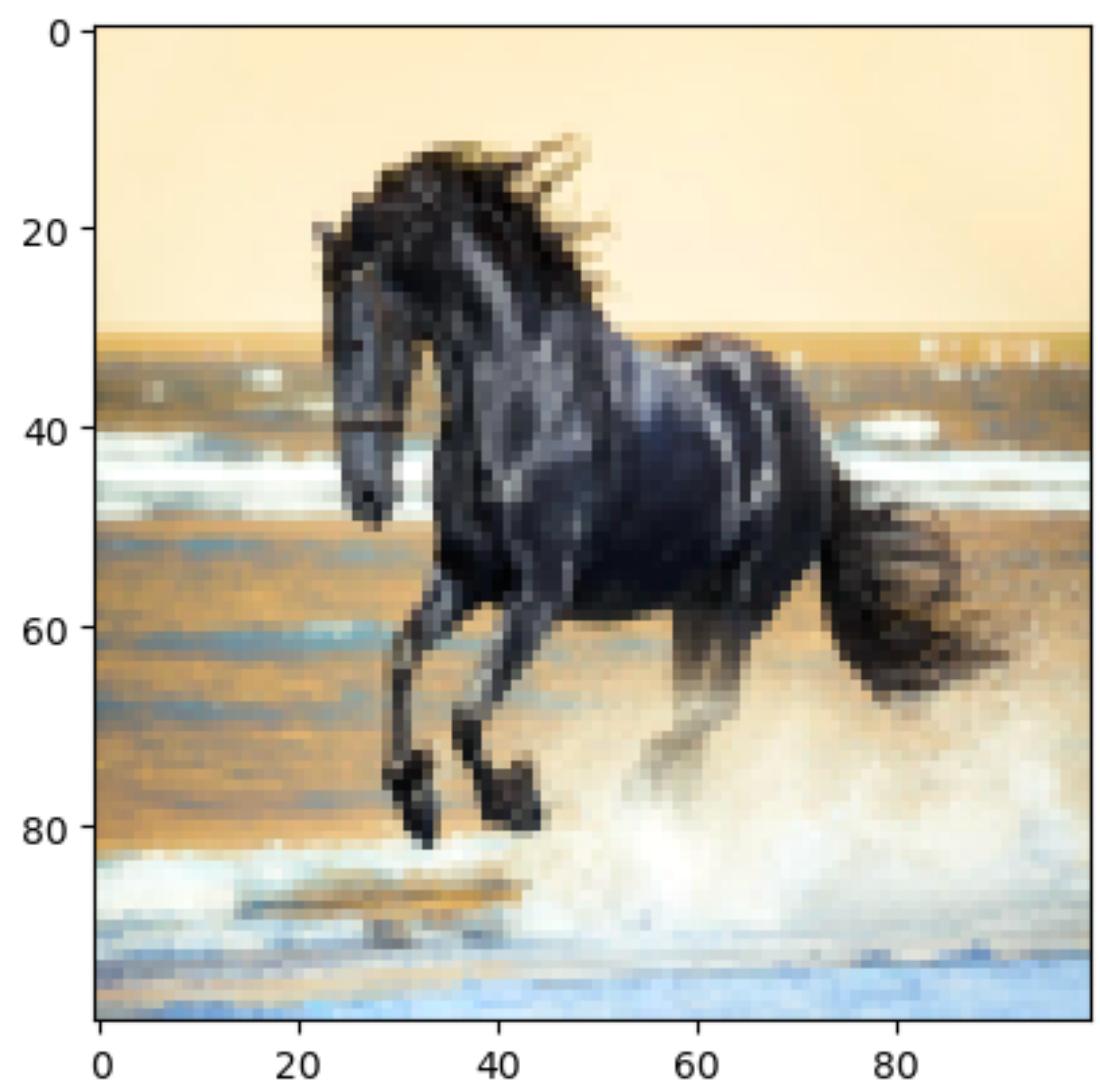
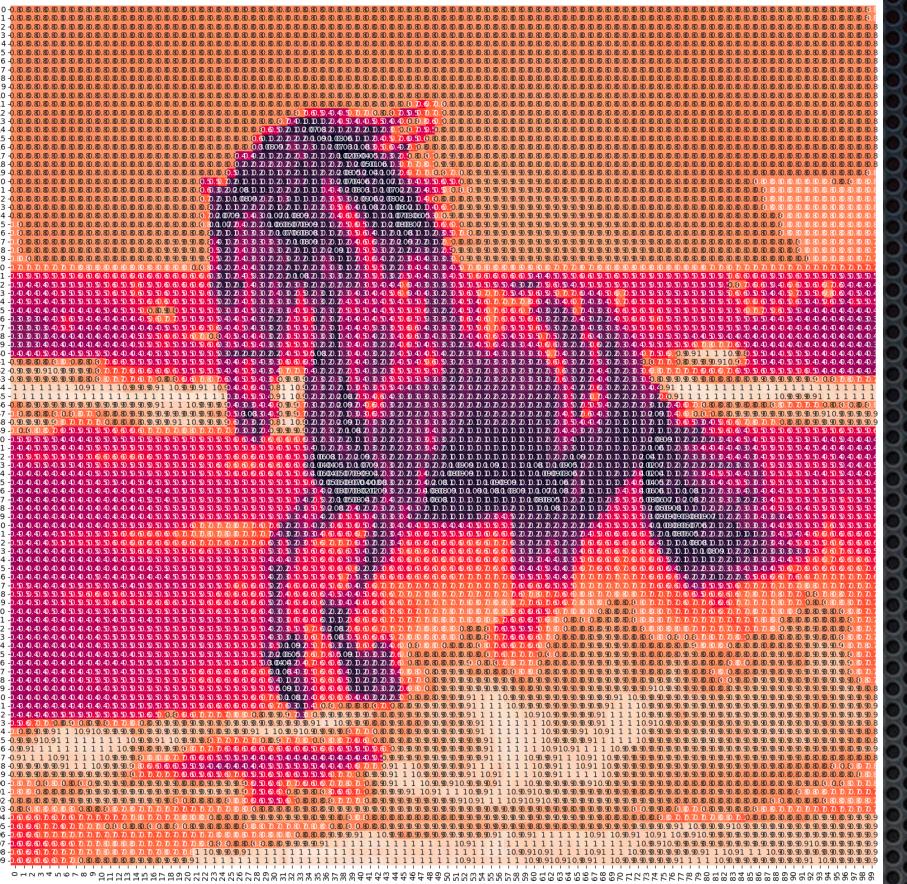
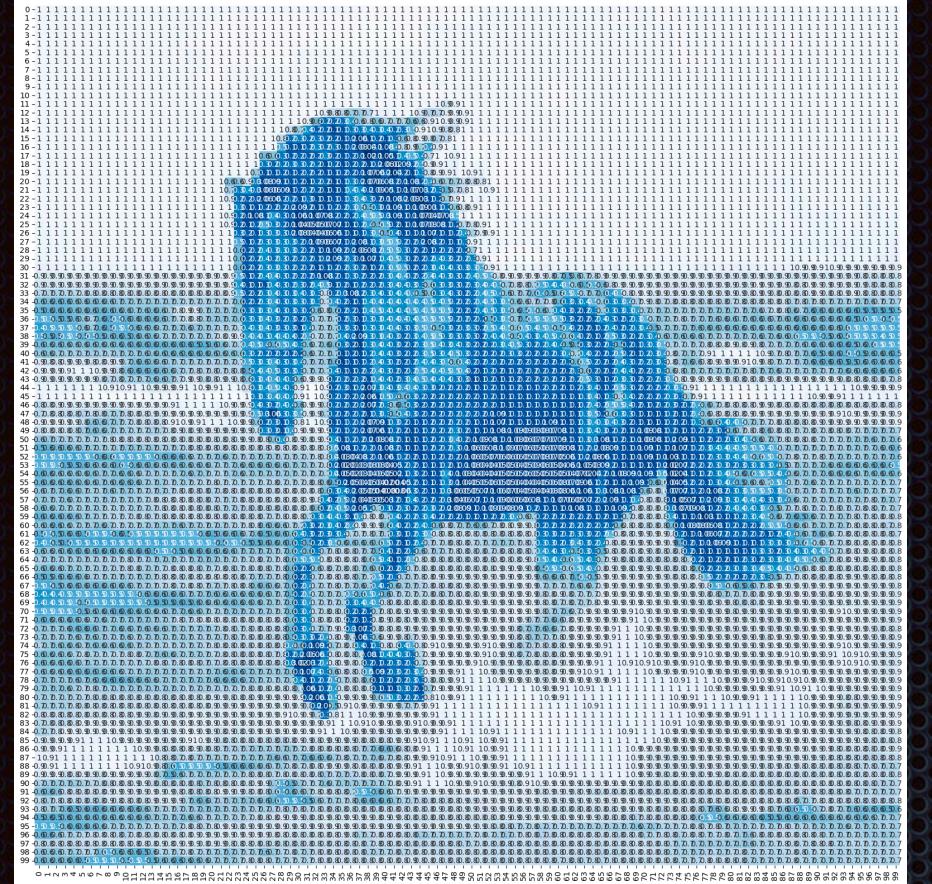
Original image
Dimension: (100,100,1)



Only binary channel

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99

Dealing with colored images



The image depth (color channels) can provide more local information than the binary image with depth 1

Interpretable AI method

Grad-CAM

Grad-Cam

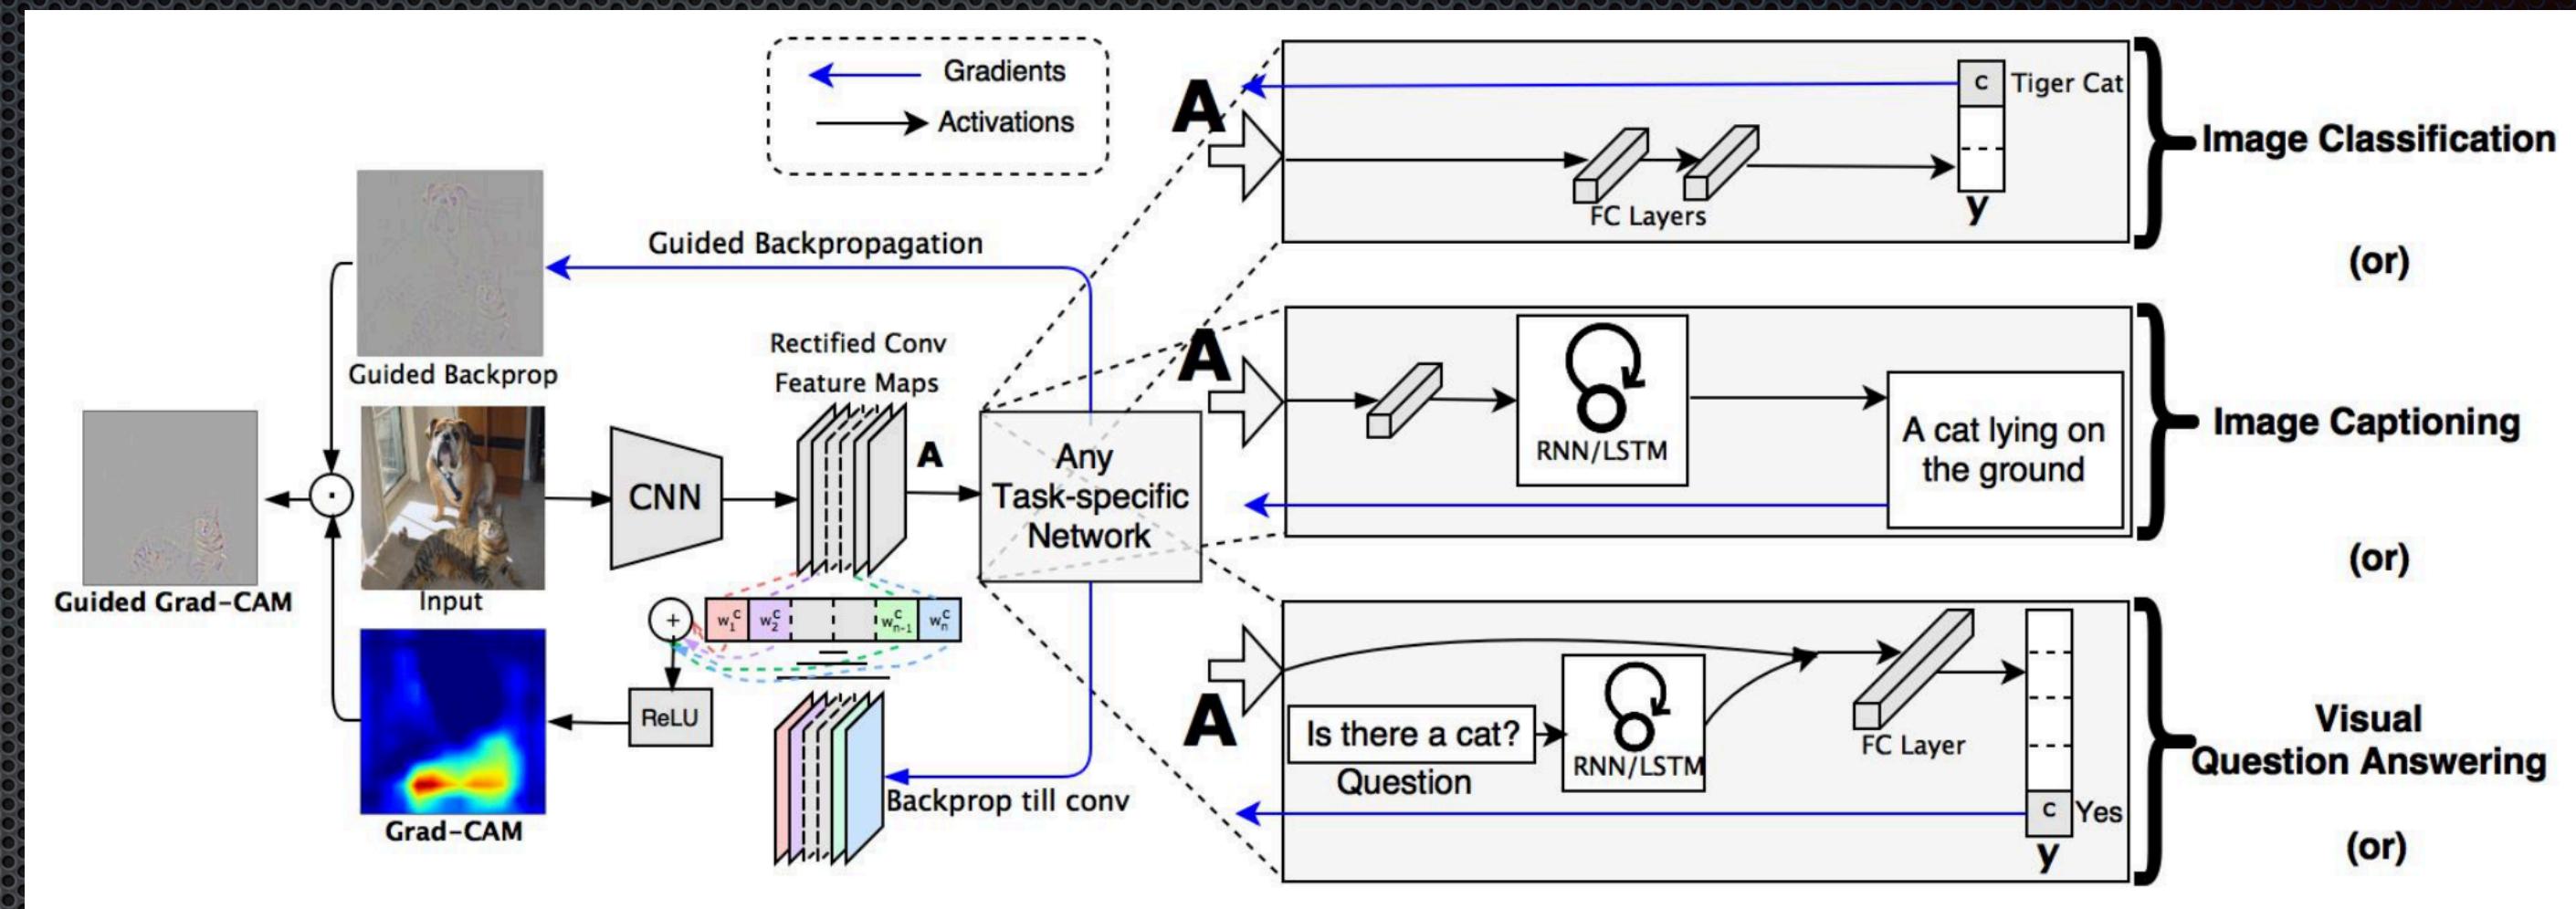
Gradient weighted Class Activation Mapping (Grad-Cam) has been first introduced in CNN model to visualize the most important pixels the model consider for his predictions.

Grad-Cam works as the following:

- After training split the model from the last convolution layer.
- Compute the output of the last convolution layer (A)
- Compute the gradient of the class score of the second half of the model
- Compute the average of the gradients with respect to the spatial coordinates

$$\alpha_k^c = \underbrace{\frac{1}{Z} \sum_i \sum_j}_{\text{global average pooling}} \underbrace{\frac{\partial y^c}{\partial A_{ij}^k}}_{\text{gradients via backprop}}$$

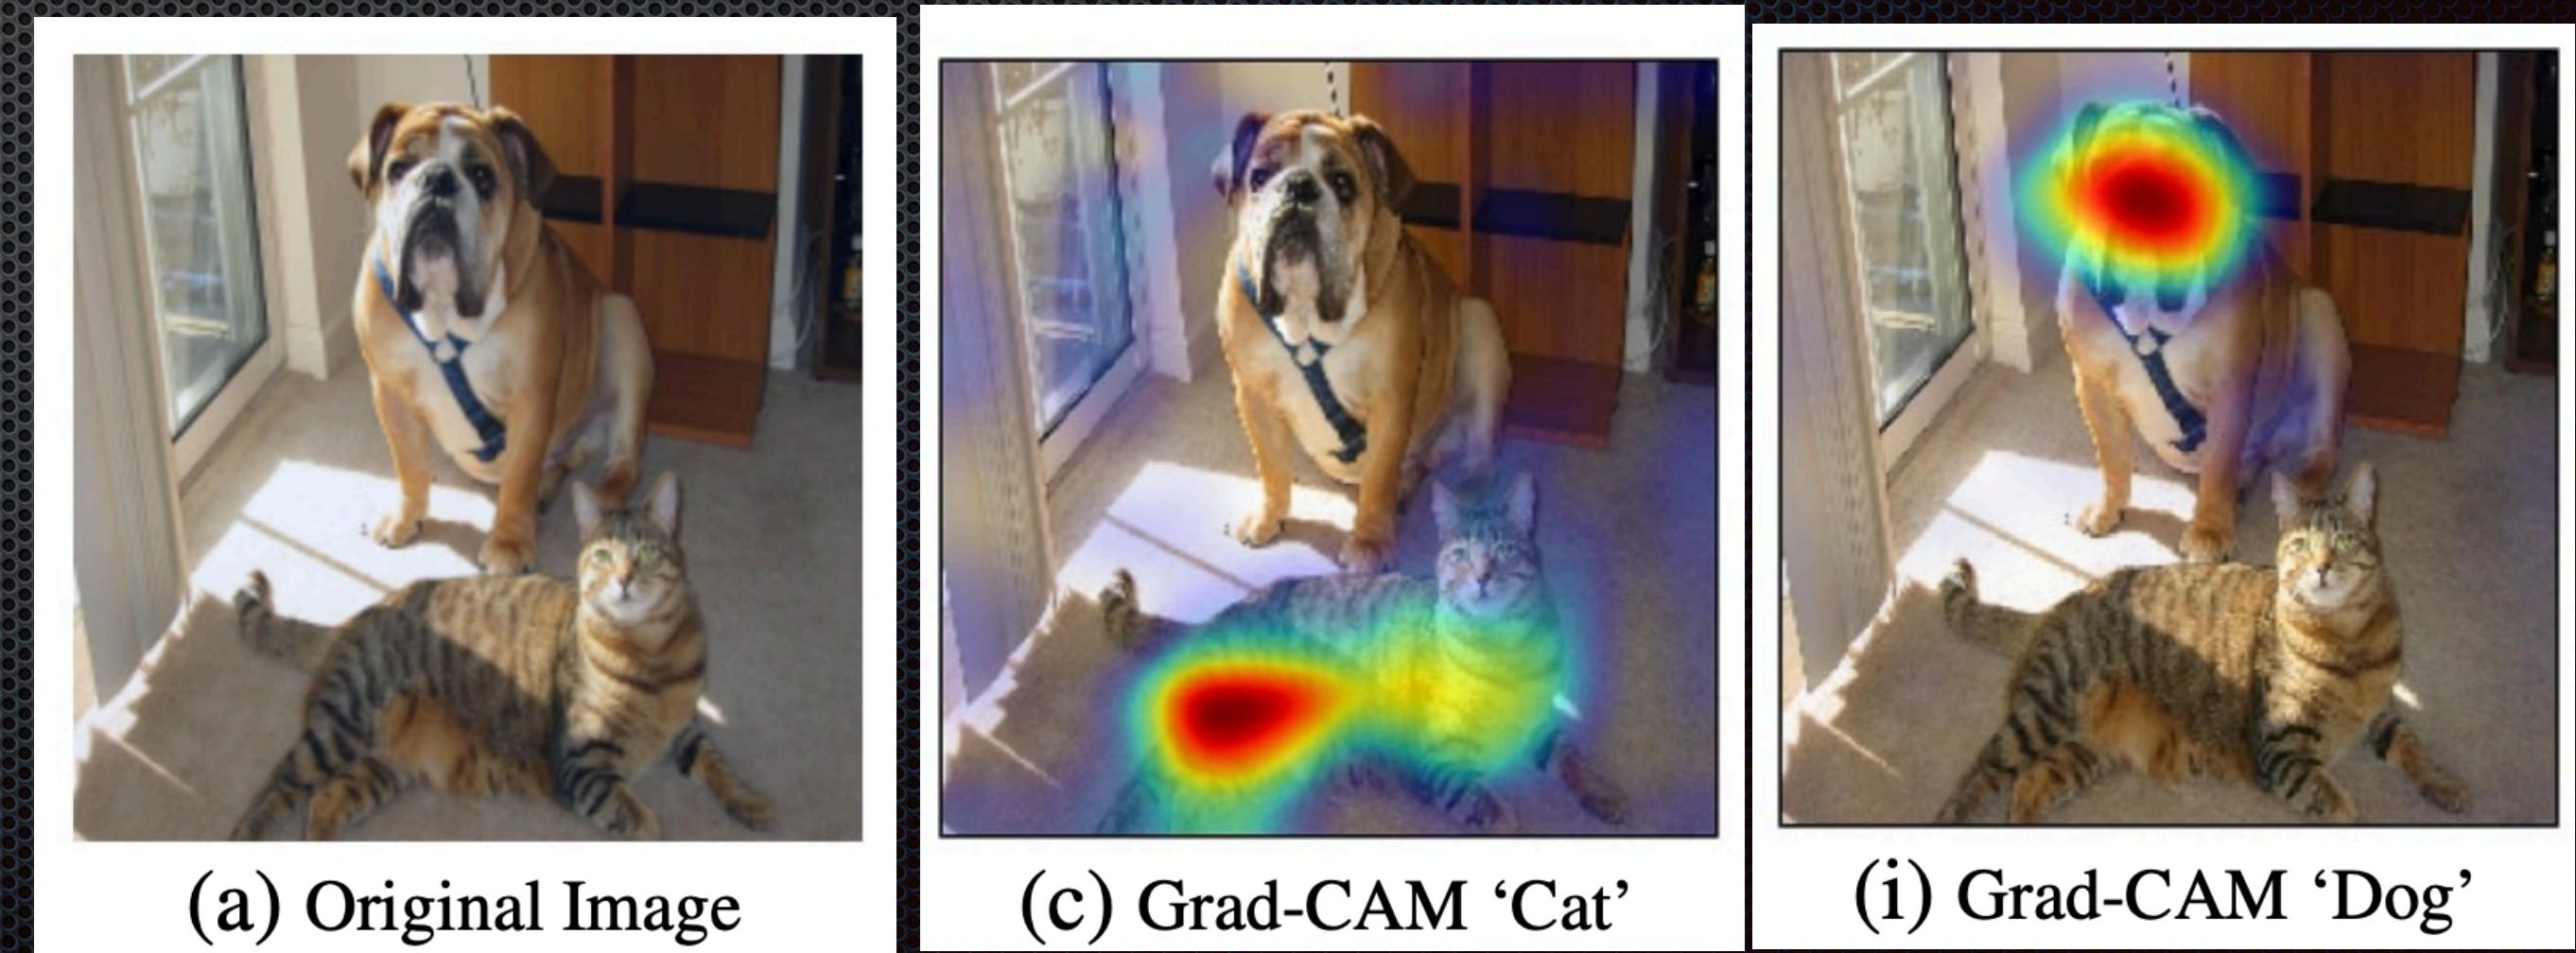
Arxiv:1610.02391



- Compute the weighted sum of the feature maps output (A)

$$L_{\text{Grad-CAM}}^c = \text{ReLU} \left(\sum_k \alpha_k^c A^k \right)$$

- The resulting heatmap indicates the spatial region in which the model focuses for predictions



Let's go coding!!